

Thread-Level Parallelism and Multithreading

- Limitations of Instruction-Level Parallelism (ILP)
- Thread-Level Parallelism and Multithreading
 - ◆ Coarse-grained Multithreading
 - ◆ Fine-grained Multithreading
 - ◆ Simultaneous Multithreading
- Case Studies of Multithreaded Processors

Limits to Instruction-Level Parallelism (ILP)

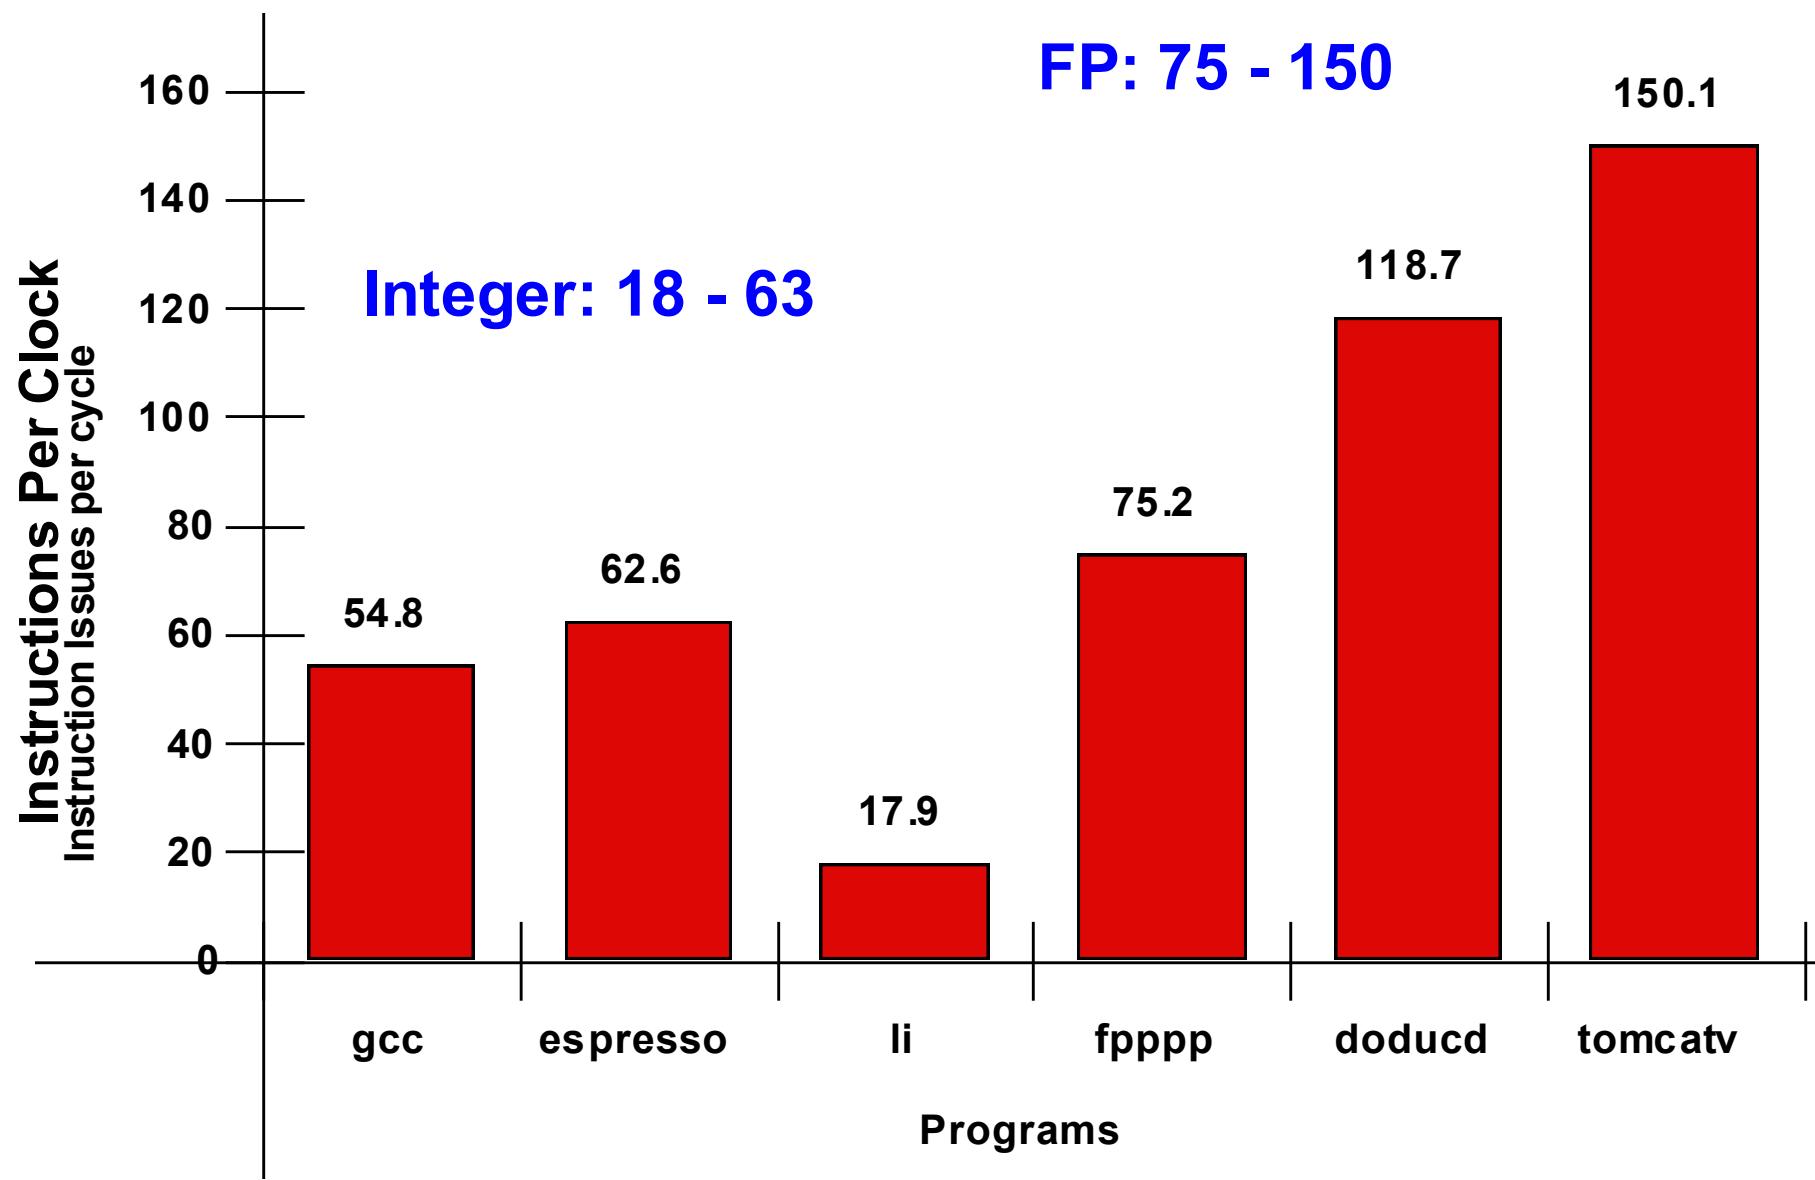
Assumptions for ideal/perfect machine to start

- Register renaming – infinite virtual registers
 - ◆ All register WAW & WAR hazards are avoided
- Memory-address alias analysis – addresses known and a load can be moved before a store provided addresses are not equal
 - ◆ The above two assumptions eliminate all but RAW
- Branch prediction – perfect; no mispredictions
- Jump prediction – all jump/returns perfectly predicted
 - ◆ no control dependencies; perfect speculation & an unbounded buffer of instructions available
- Perfect caches – 1 cycle latency for all instructions
- Unlimited Issue – Unlimited instructions issued per clock cycle

Limits to ILP HW Model Comparison

| | Perfect Model | Power 5 |
|--------------------------------------|---------------|---|
| Instructions Issued per Clock | Infinite | 4 |
| Instruction Window Size | Infinite | 200 |
| Renaming Registers | Infinite | 48 integer + 40 Floating-point |
| Branch Prediction | Perfect | 2% to 6% misprediction (Tournament Branch Predictor) |
| Cache | Perfect | 64KI, 32KD, 1.92M L2, 36M L3 |
| Memory Alias Analysis | Perfect | ?? |

Upper Limit to ILP: Ideal Machine



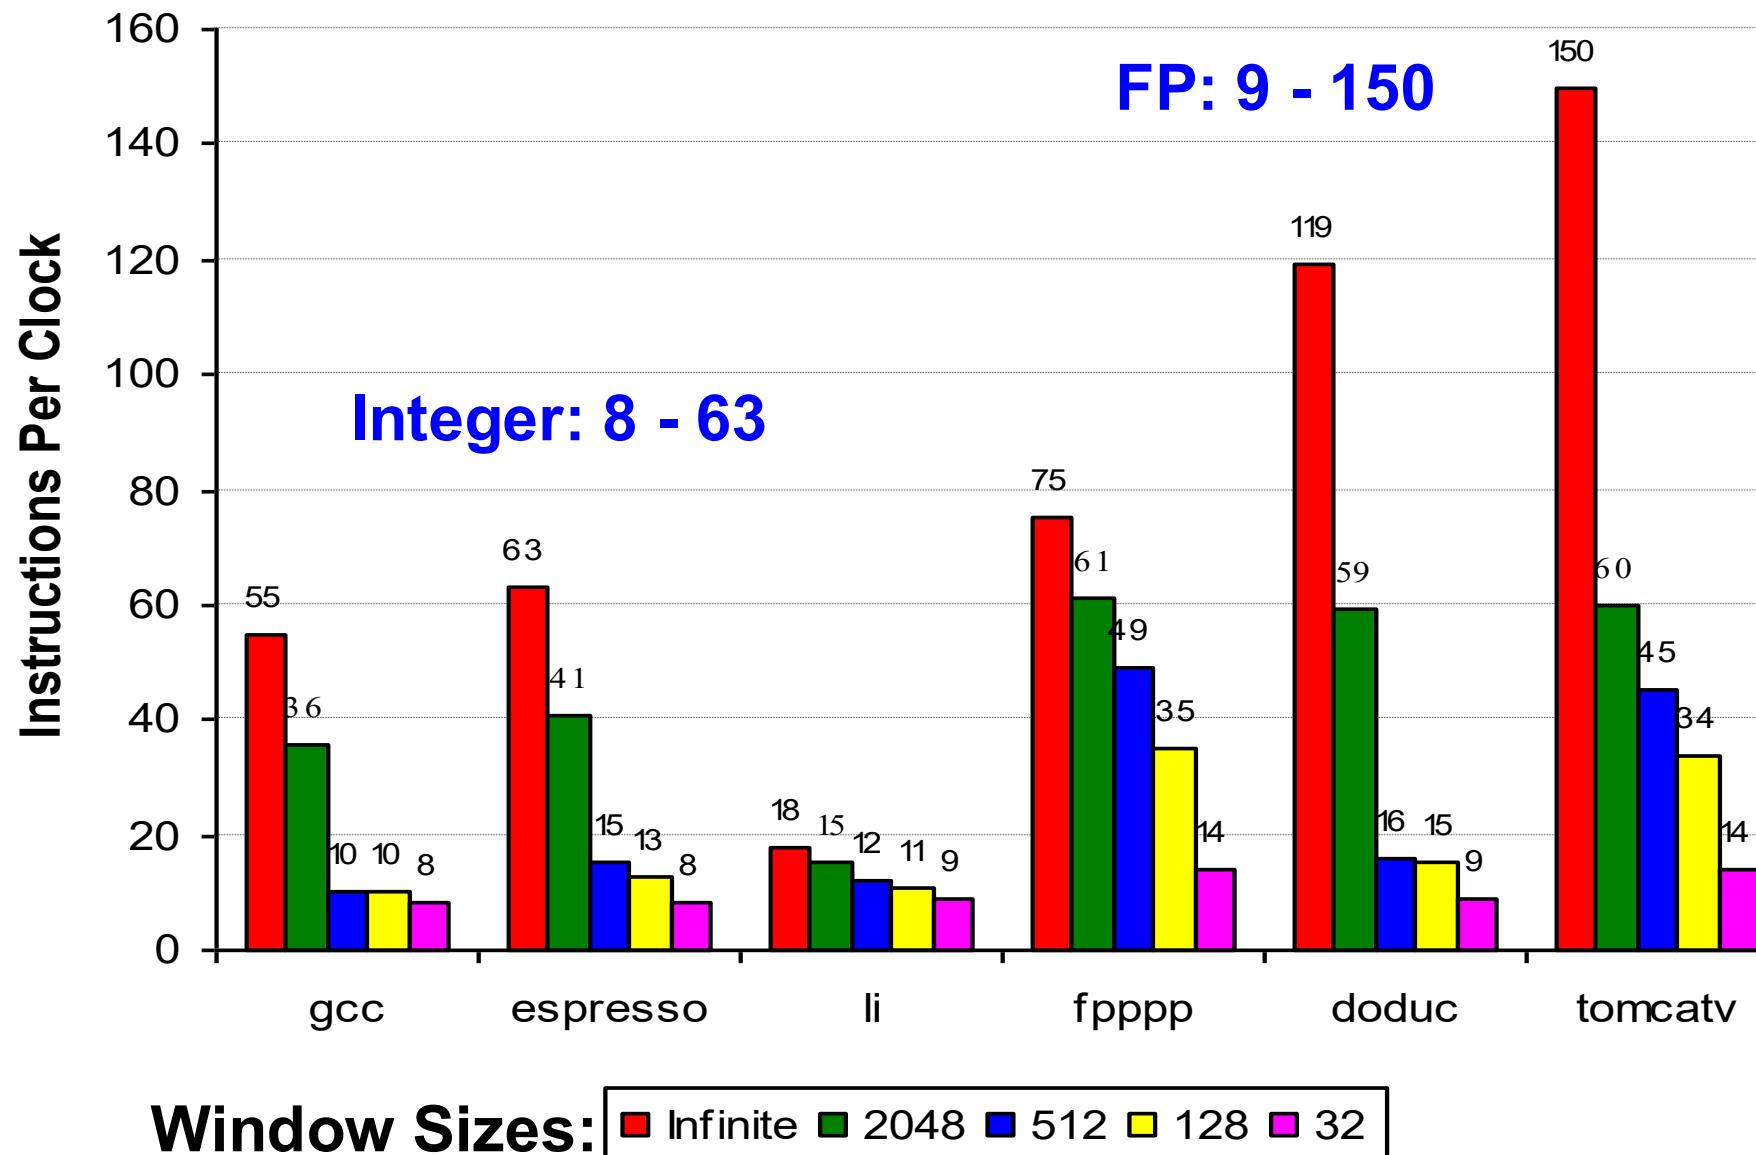
Practical Check: Maximum Issue Count

- How close a real dynamically scheduled, speculative processor come to the ideal one?
 - ◆ Look arbitrarily far ahead predicting all branches
 - ◆ Rename all register uses to avoid WAR/WAW
 - ◆ Determine data dependencies
 - ◆ Determine memory dependencies
 - ◆ Enough parallel units
- Determine if n issuing instructions have any dependencies
 - ◆ RAW: $2n - 2 + 2n - 4 + \dots + 2 = 2 \sum_{i=1}^{n-1} i = 2.n(n-1)/2 = n^2 - n$
 - ◆ WAR: $n^2 - n$
 - ◆ WAW: $(n-1) + (n-2) + \dots + 1 = n(n-1)/2 = (n^2 - n)/2$
 - ◆ Issuing 50 instructions require 6125 ($2450 + 2450 + 1225$) comparisons

Limits to ILP HW Model Comparison

| | New Model | Perfect Model | Power 5 |
|-------------------------------|----------------------------|---------------|---|
| Instructions Issued per Clock | Infinite | Infinite | 4 |
| Instruction Window Size | Infinite, 2K, 512, 128, 32 | Infinite | 200 |
| Renaming Registers | Infinite | Infinite | 48 integer + 40 Floating-point |
| Branch Prediction | Perfect | Perfect | 2% to 6% misprediction (Tournament Branch Predictor) |
| Cache | Perfect | Perfect | 64KI, 32KD, 1.92B L2, 36M L3 |
| Memory Alias | Perfect | Perfect | ?? |

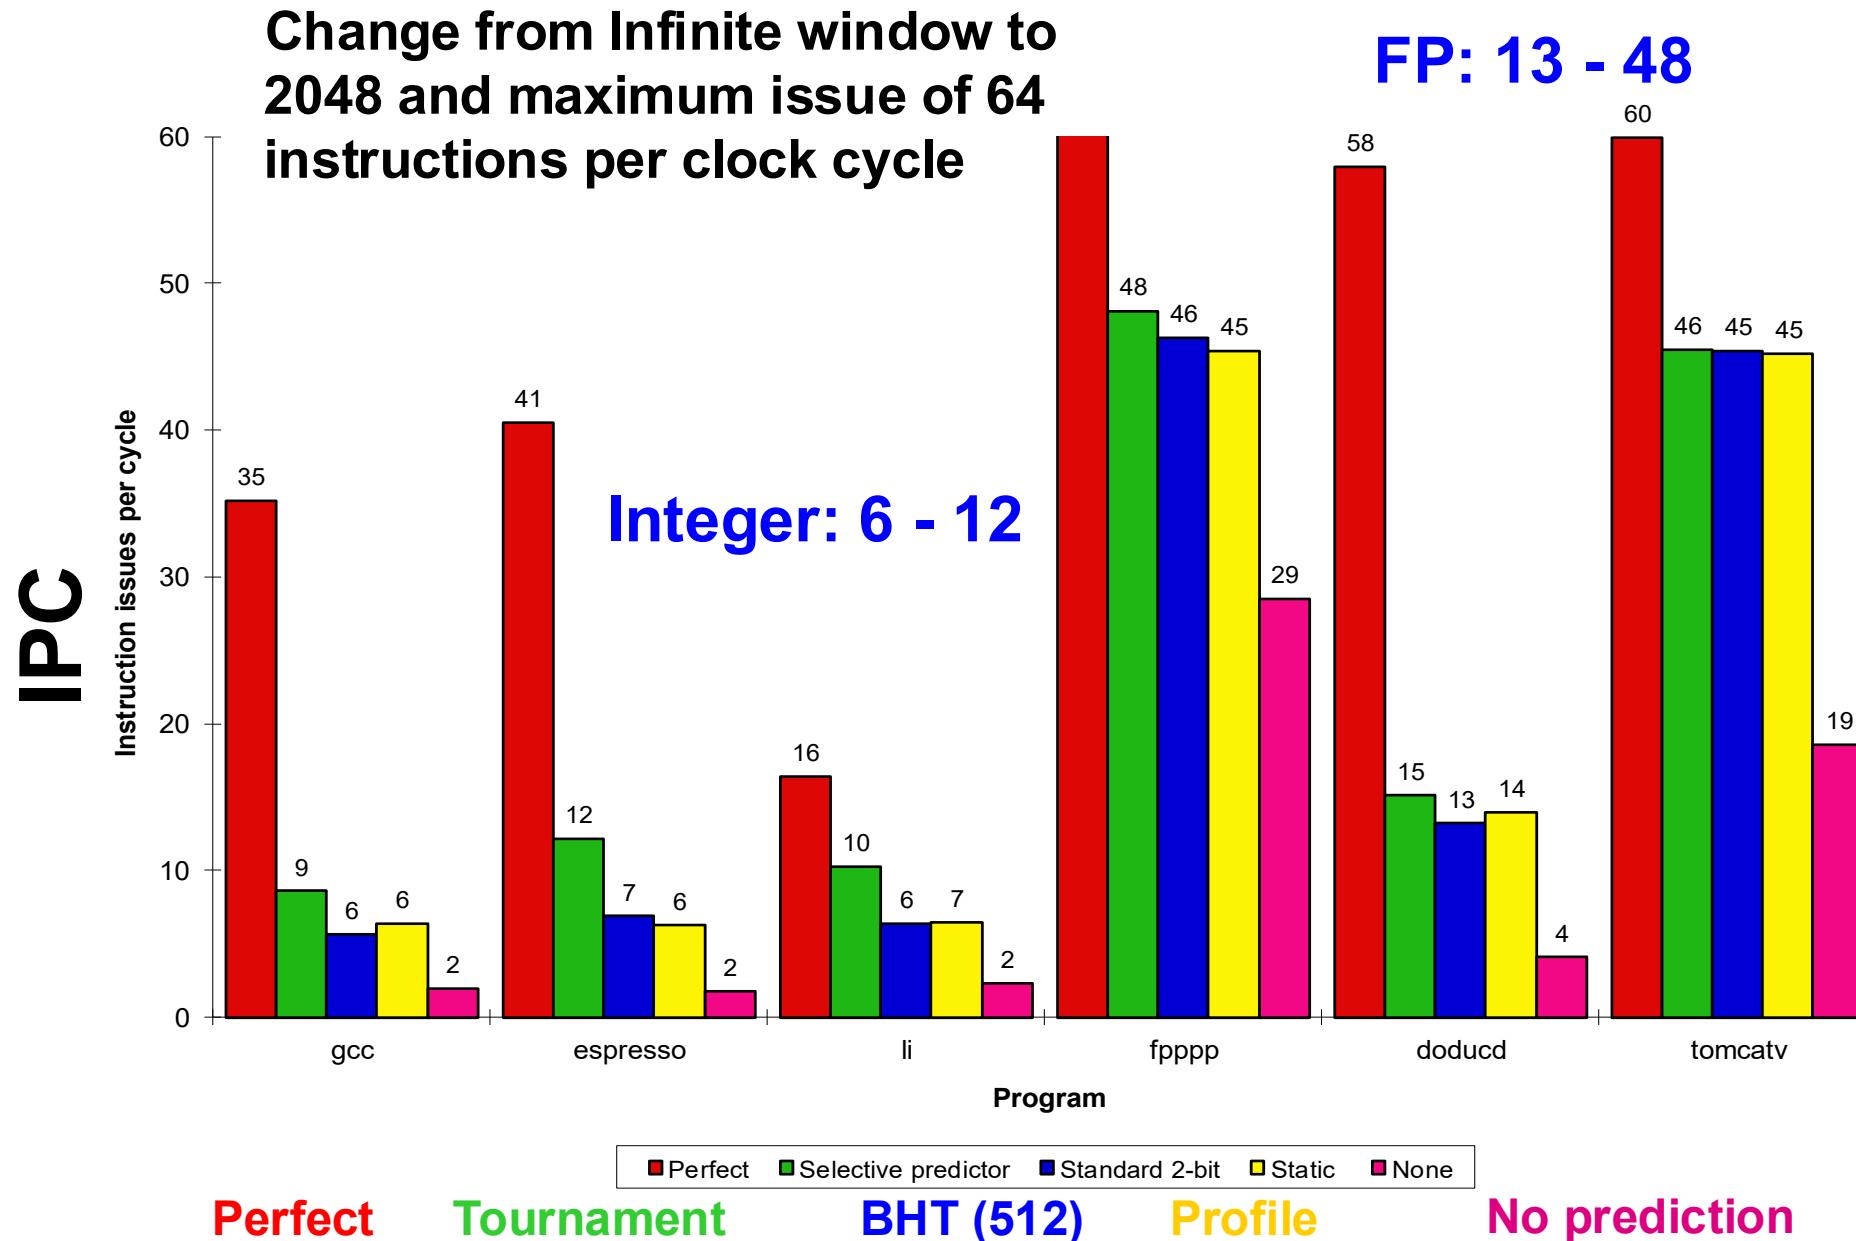
More Realistic HW: Window Impact



Limits to ILP HW Model Comparison

| | New Model | Perfect Model | Power 5 |
|-------------------------------|--|---------------|---|
| Instructions Issued per Clock | 64 | Infinite | 4 |
| Instruction Window Size | 2048 | Infinite | 200 |
| Renaming Registers | Infinite | Infinite | 48 integer + 40 Floating-point |
| Branch Prediction | Perfect, 8K Tournament, 512 2-bit, profile, none | Perfect | 2% to 6% misprediction (Tournament Branch Predictor) |
| Cache | Perfect | Perfect | 64KI, 32KD, 1.92B L2, 36M L3 |
| Memory Alias | Perfect | Perfect | ?? |

More Realistic HW: Branch Impact

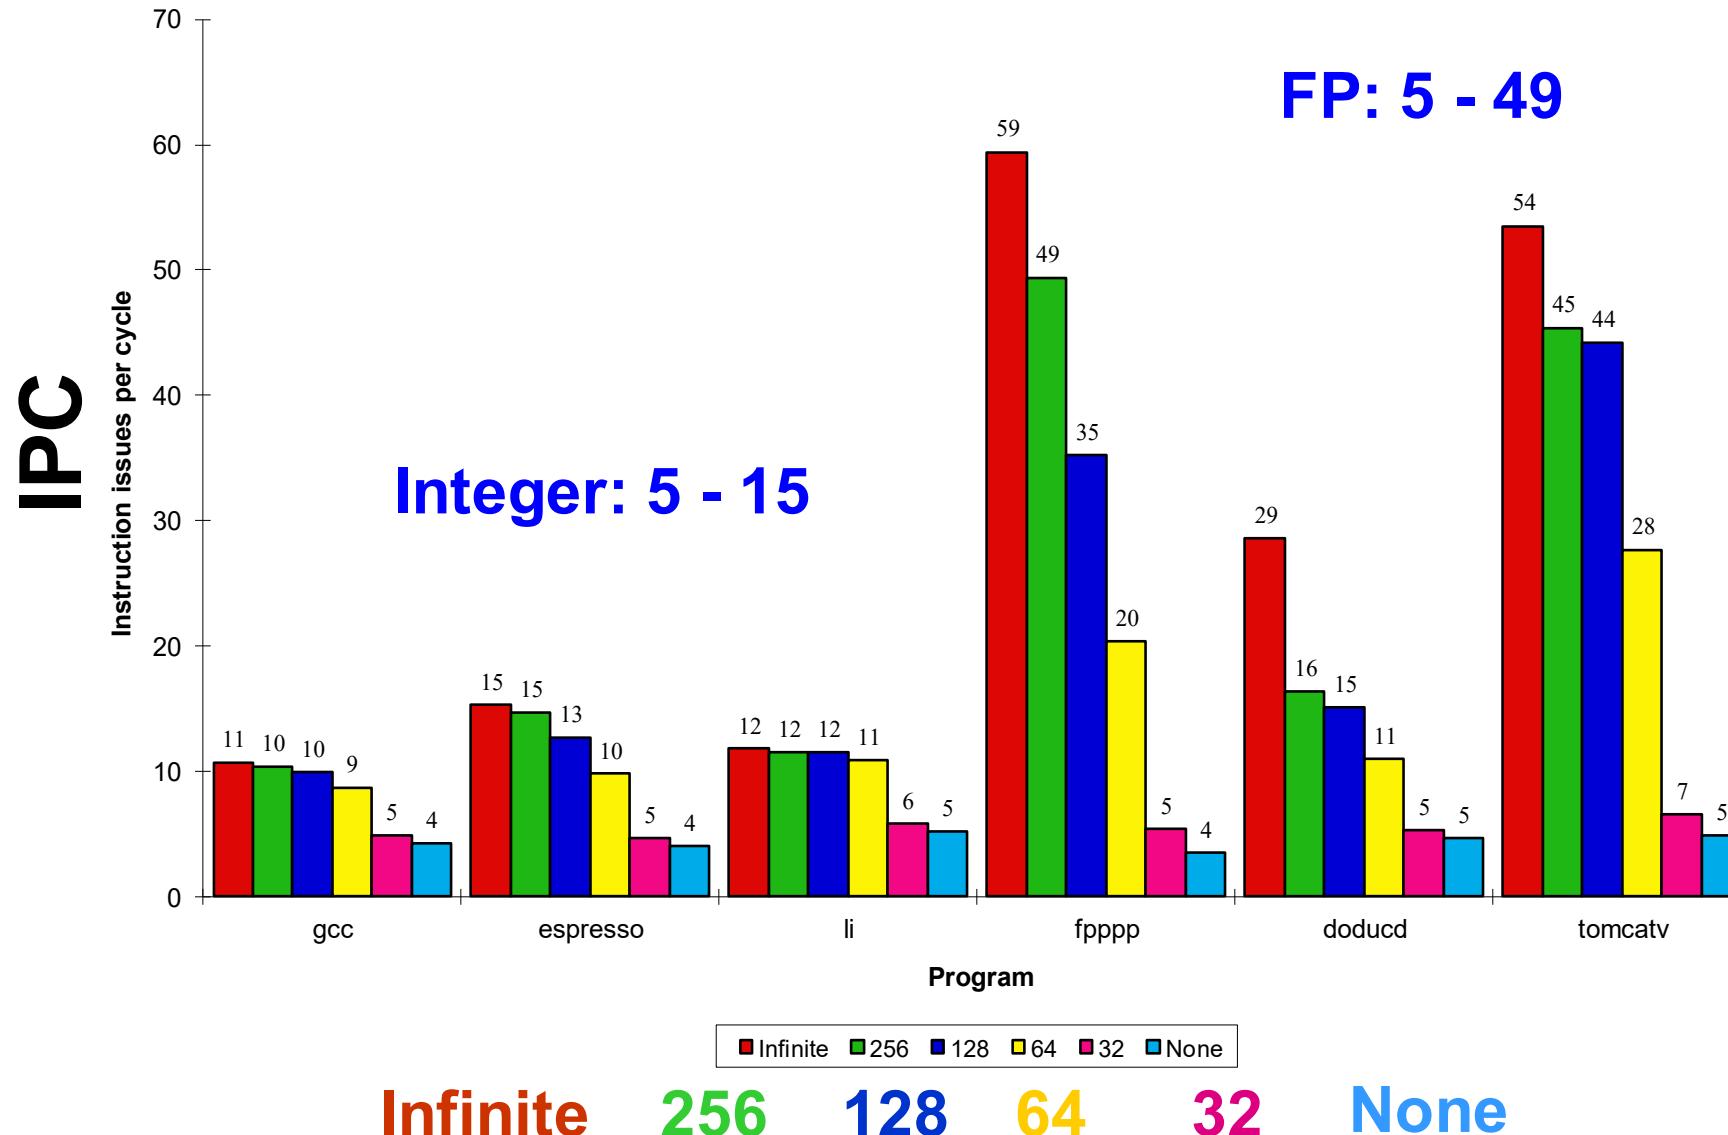


Limits to ILP HW Model Comparison

| | New Model | Perfect Model | Power 5 |
|-------------------------------|---|---------------|---|
| Instructions Issued per Clock | 64 | Infinite | 4 |
| Instruction Window Size | 2048 | Infinite | 200 |
| Renaming Registers | Infinite, 256, 128, 64, 32, none | Infinite | 48 integer + 40 Floating-point |
| Branch Prediction | 8K 2-bit | Perfect | 2% to 6% misprediction (Tournament Branch Predictor) |
| Cache | Perfect | Perfect | 64KI, 32KD, 1.92B L2, 36M L3 |
| Memory Alias | Perfect | Perfect | ?? |

Renaming Register Impact (N int + N fp)

Change 2048 instr window, 64 instr issue, 8K 2 level prediction

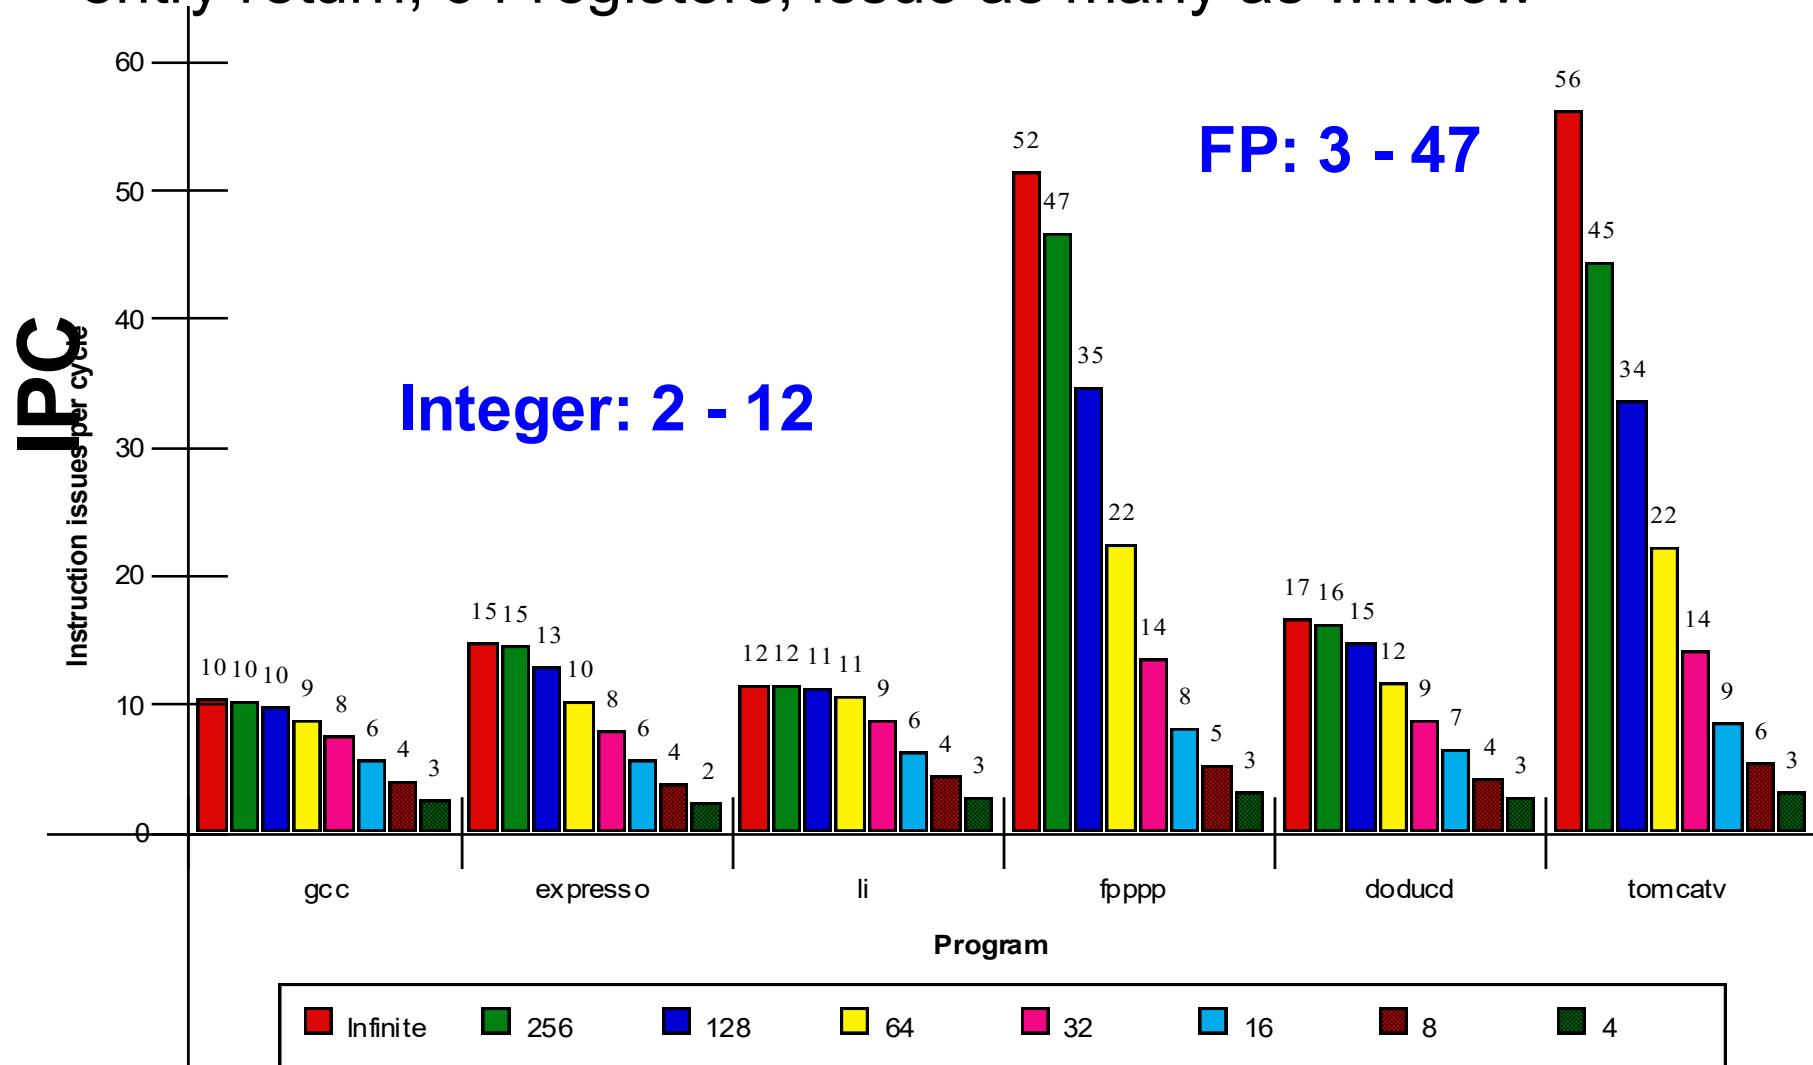


Limits to ILP HW Model Comparison

| | New Model | Perfect Model | Power 5 |
|-------------------------------|----------------------------|---------------|---|
| Instructions Issued per Clock | 64 | Infinite | 4 |
| Instruction Window Size | Infinite, 256, 128, 64, 32 | Infinite | 200 |
| Renaming Registers | 64 Int + 64 FP | Infinite | 48 integer + 40 Floating-point |
| Branch Prediction | 1K 2-bit | Perfect | 2% to 6% misprediction (Tournament Branch Predictor) |
| Cache | Perfect | Perfect | 64KI, 32KD, 1.92B L2, 36M L3 |
| Memory Alias | HW dis-ambiguation | Perfect | ?? |

Realistic HW: Window Impact

Perfect disambiguation (HW), 1K Selective Prediction, 16 entry return, 64 registers, issue as many as window



How to Exceed ILP Limits?

- These are not laws of physics; just practical limits for today, and perhaps overcome via research
- Compiler & ISA advances can change results
- WAR and WAW hazards through memory: eliminated WAW and WAR hazards through register renaming, but not in memory usage
 - ◆ Can get conflicts via allocation of stack frames as a called procedure reuses the memory addresses of a previous frame on the stack
- Unnecessary dependencies e.g., loop variable can be removed by loop unrolling
- The data flow limit can be improved by value prediction

Thread-Level Parallelism and Multithreading

- Limitations of Instruction-Level Parallelism (ILP)
- Thread-Level Parallelism and Multithreading
 - ◆ Coarse-grained Multithreading
 - ◆ Fine-grained Multithreading
 - ◆ Simultaneous Multithreading
- Case Studies of Multithreaded Processors

Performance beyond Single Thread ILP

- There can be higher natural parallelism in some applications
 - ◆ Database or scientific applications
 - ◆ Explicit Data Level Parallelism or Thread Level Parallelism
- Data-Level Parallelism: Perform identical operations on data
 - ◆ We will discuss this later.
- Thread-Level Parallelism
 - ◆ Thread is a process with own instructions and data
 - ◆ Thread may be a process part of a parallel program of multiple processes, or it may be an independent program
 - ◆ Each thread has all the state (instructions, data, PC, register state, and so on) necessary to allow it to execute

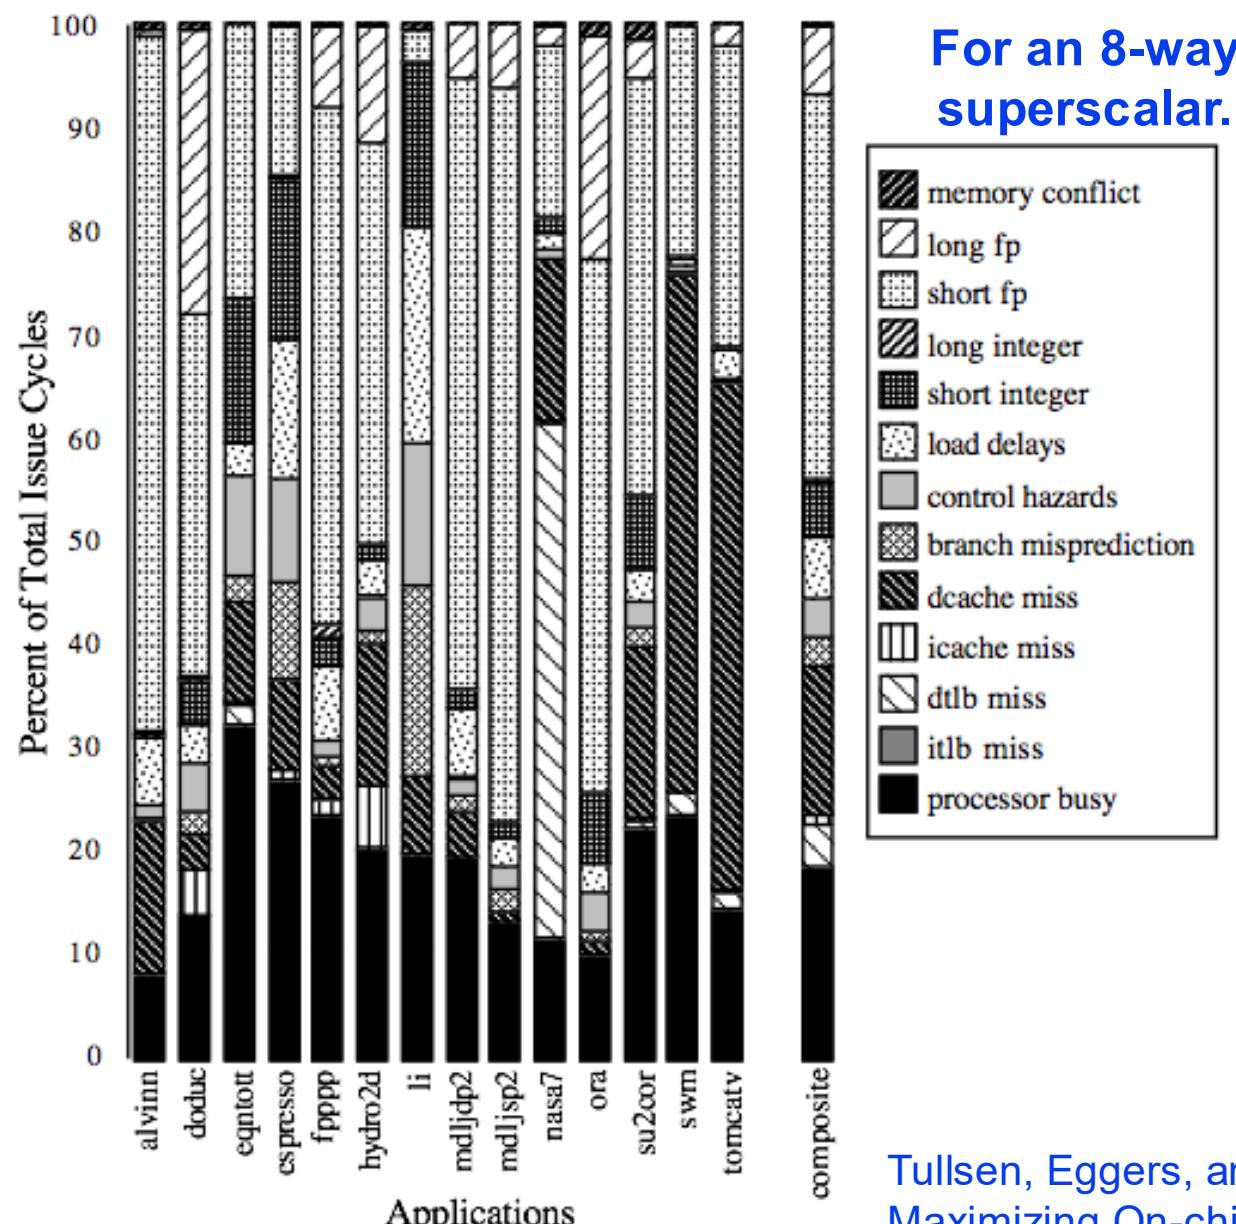
Thread Level Parallelism (TLP)

- ILP exploits implicit parallel operations within a loop or straight-line code segment
- TLP are represented using multiple threads of execution that are inherently parallel
- Goal: Use multiple instruction streams to improve
 1. Throughput of computers that run many programs
 2. Execution time of multi-threaded programs
- TLP is more cost-effective to exploit than ILP

Do both ILP and TLP?

- Functional units are often idle in data path designed for ILP because of either stalls or dependences
- TLP and ILP exploit two different kinds of parallel structure in a program
- Can a processor oriented at ILP exploit TLP?
- Could the TLP be used as a source of independent instructions that might keep the processor busy during stalls?
- Could TLP be used to employ the functional units that would otherwise lie idle when insufficient ILP exists?

For Most Applications, Most Execution Units are Idle

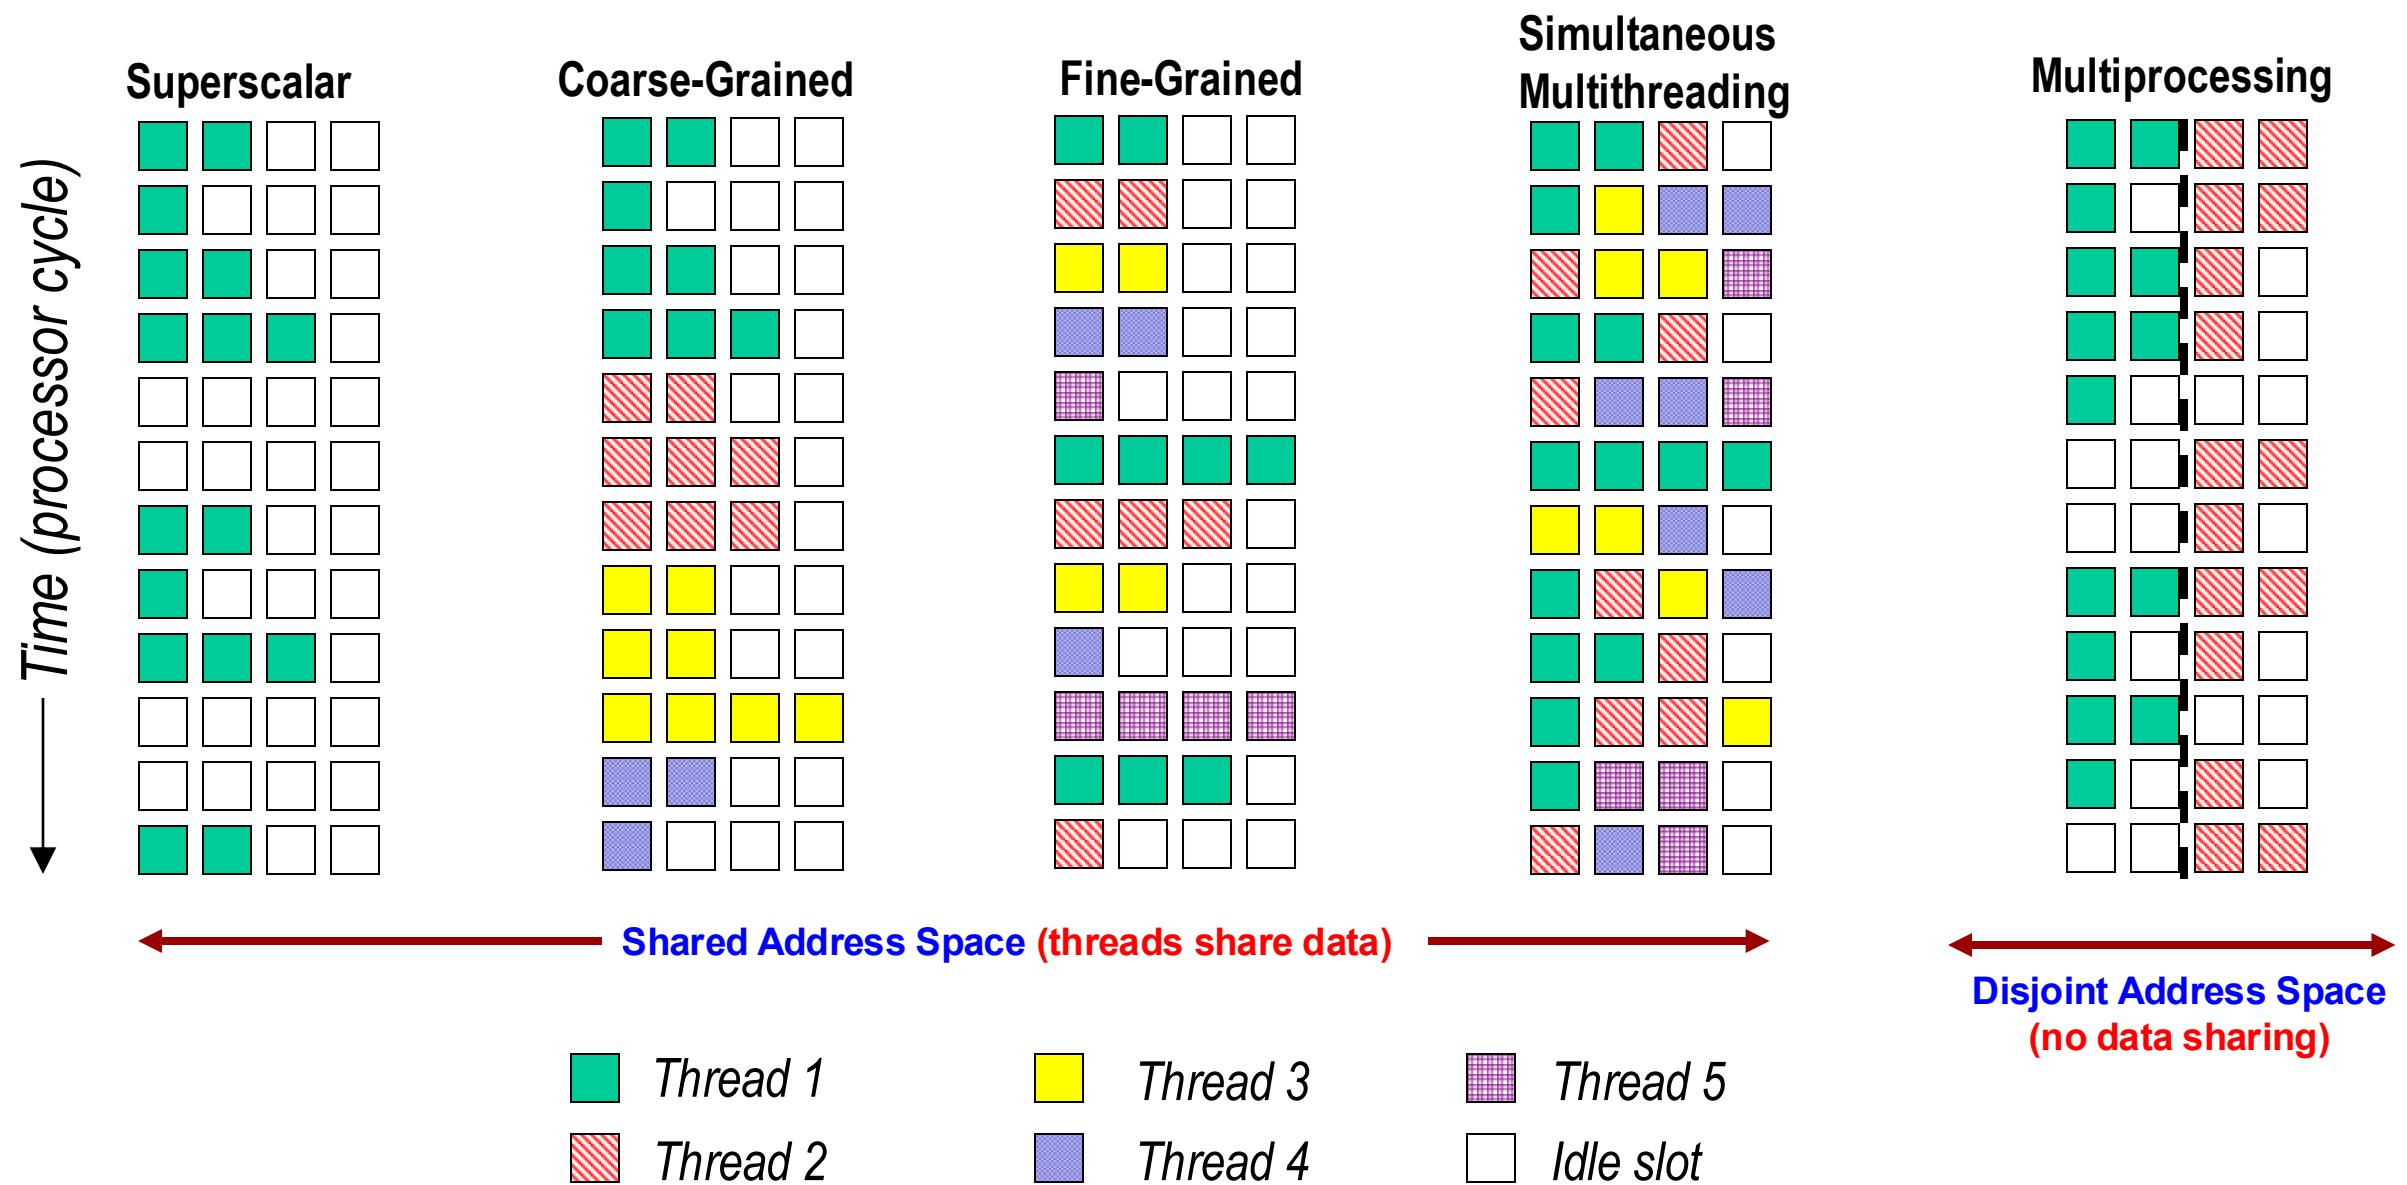


Tullsen, Eggers, and Levy, "Simultaneous Multithreading: Maximizing On-chip Parallelism", ISCA 1995.

New Approach: Multithreaded Execution

- **Multithreading:** multiple threads to share the functional units of one processor via overlapping
 - ◆ processor must duplicate independent state of each thread e.g., a separate copy of register file, a separate PC, and for running independent programs, a separate page table
 - ◆ memory shared through the virtual memory mechanisms, which already support multiple processes
 - ◆ HW for fast thread switch; much faster than full process switch \approx 100s to 1000s of clocks
- When to switch?
 - ◆ Alternate instruction per thread (fine grain)
 - ◆ When a thread is stalled, perhaps for a cache miss, another thread can be executed (coarse grain)

Multithreaded Categories



Coarse-Grained Multithreading

- Switches threads only on costly stalls, such as L2 cache miss
- Advantages
 - ◆ Relieves need to have very fast thread-switching
 - ◆ Doesn't slow down thread, since instructions from other threads issued only when the thread reaches a costly stall
- Disadvantage is hard to overcome throughput losses from shorter stalls, due to pipeline start-up costs
 - ◆ Since CPU issues instructions from 1 thread, when a stall occurs, the pipeline must be emptied or frozen
 - ◆ New thread must fill pipeline before instructions can complete
- Because of this start-up overhead, coarse-grained multithreading is better for reducing penalty of high cost stalls, where pipeline refill << stall time
- Used in IBM AS/400

Fine-Grained Multithreading

- Switches between threads on each instruction, causing the execution of multiple threads to be interleaved
 - ◆ Usually done in a round-robin fashion, skipping any stalled threads
- CPU must be able to switch threads each clock
- **Advantage** – it can hide both short and long stalls, since instructions from other threads execute when one thread stalls
- **Disadvantage** – it slows down execution of individual threads, since a thread ready to execute without stalls will be delayed by instructions from other threads
- Used in Sun's Niagara (will see later)

Simultaneous Multithreading (SMT)

- Dynamically scheduled processor already has many HW mechanisms to support multithreading
 - ◆ Large set of virtual registers that can be used to hold the register sets of independent threads
 - ◆ Register renaming provides unique register identifiers, so instructions from multiple threads can be mixed in datapath without confusing sources and destinations across threads
 - ◆ Out-of-order completion allows the threads to execute out of order, and get better utilization of the hardware
- Add a per thread renaming table and keeping separate PCs
 - ◆ Independent commitment can be supported by logically keeping a separate reorder buffer for each thread

Simultaneous Multi-Threading (SMT)

One thread, 8 units

Cycle M M FX FX FP FP BR CC

| | | | | | | | | | |
|---|---|---|---|---|---|--|---|--|---|
| 1 | Y | | | | | | | | Y |
| 2 | | Y | | | | | Y | | |
| 3 | | | Y | Y | | | | | |
| 4 | | | | | | | | | |
| 5 | | | | | | | | | |
| 6 | | | | | | | | | |
| 7 | Y | | | Y | | | Y | | |
| 8 | | Y | | | Y | | | | |
| 9 | | | Y | | | | | | |

Two threads, 8 units

Cycle M M FX FX FP FP BR CC

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | Y | | B | B | | | | | Y |
| 2 | | Y | | B | | | | B | Y |
| 3 | B | | | | | Y | Y | | |
| 4 | B | | | | | | | B | |
| 5 | | B | | | | | | | B |
| 6 | | | | | | | | | |
| 7 | Y | | | B | Y | B | Y | | |
| 8 | | Y | | | B | Y | | B | |
| 9 | B | B | | | Y | | Y | B | |

M = Load/Store, FX = Fixed Point, FP = Floating Point, BR = Branch, CC = Condition Codes

Design Challenges in SMT

- Since SMT makes sense only with fine-grained implementation, impact of fine-grained scheduling on single thread performance.
 - ◆ A preferred thread approach sacrifices neither throughput nor single-thread performance.
 - ◆ Unfortunately, with a preferred thread, the processor is likely to sacrifice some throughput, when preferred thread stalls
- Larger register file needed to hold multiple contexts
- Not affecting clock cycle time, especially in
 - ◆ Instruction issue - more candidate instructions need to be considered
 - ◆ Instruction completion - choosing which instructions to commit may be challenging
- Ensuring that cache and TLB conflicts generated by SMT do not degrade performance

Thread-Level Parallelism and Multithreading

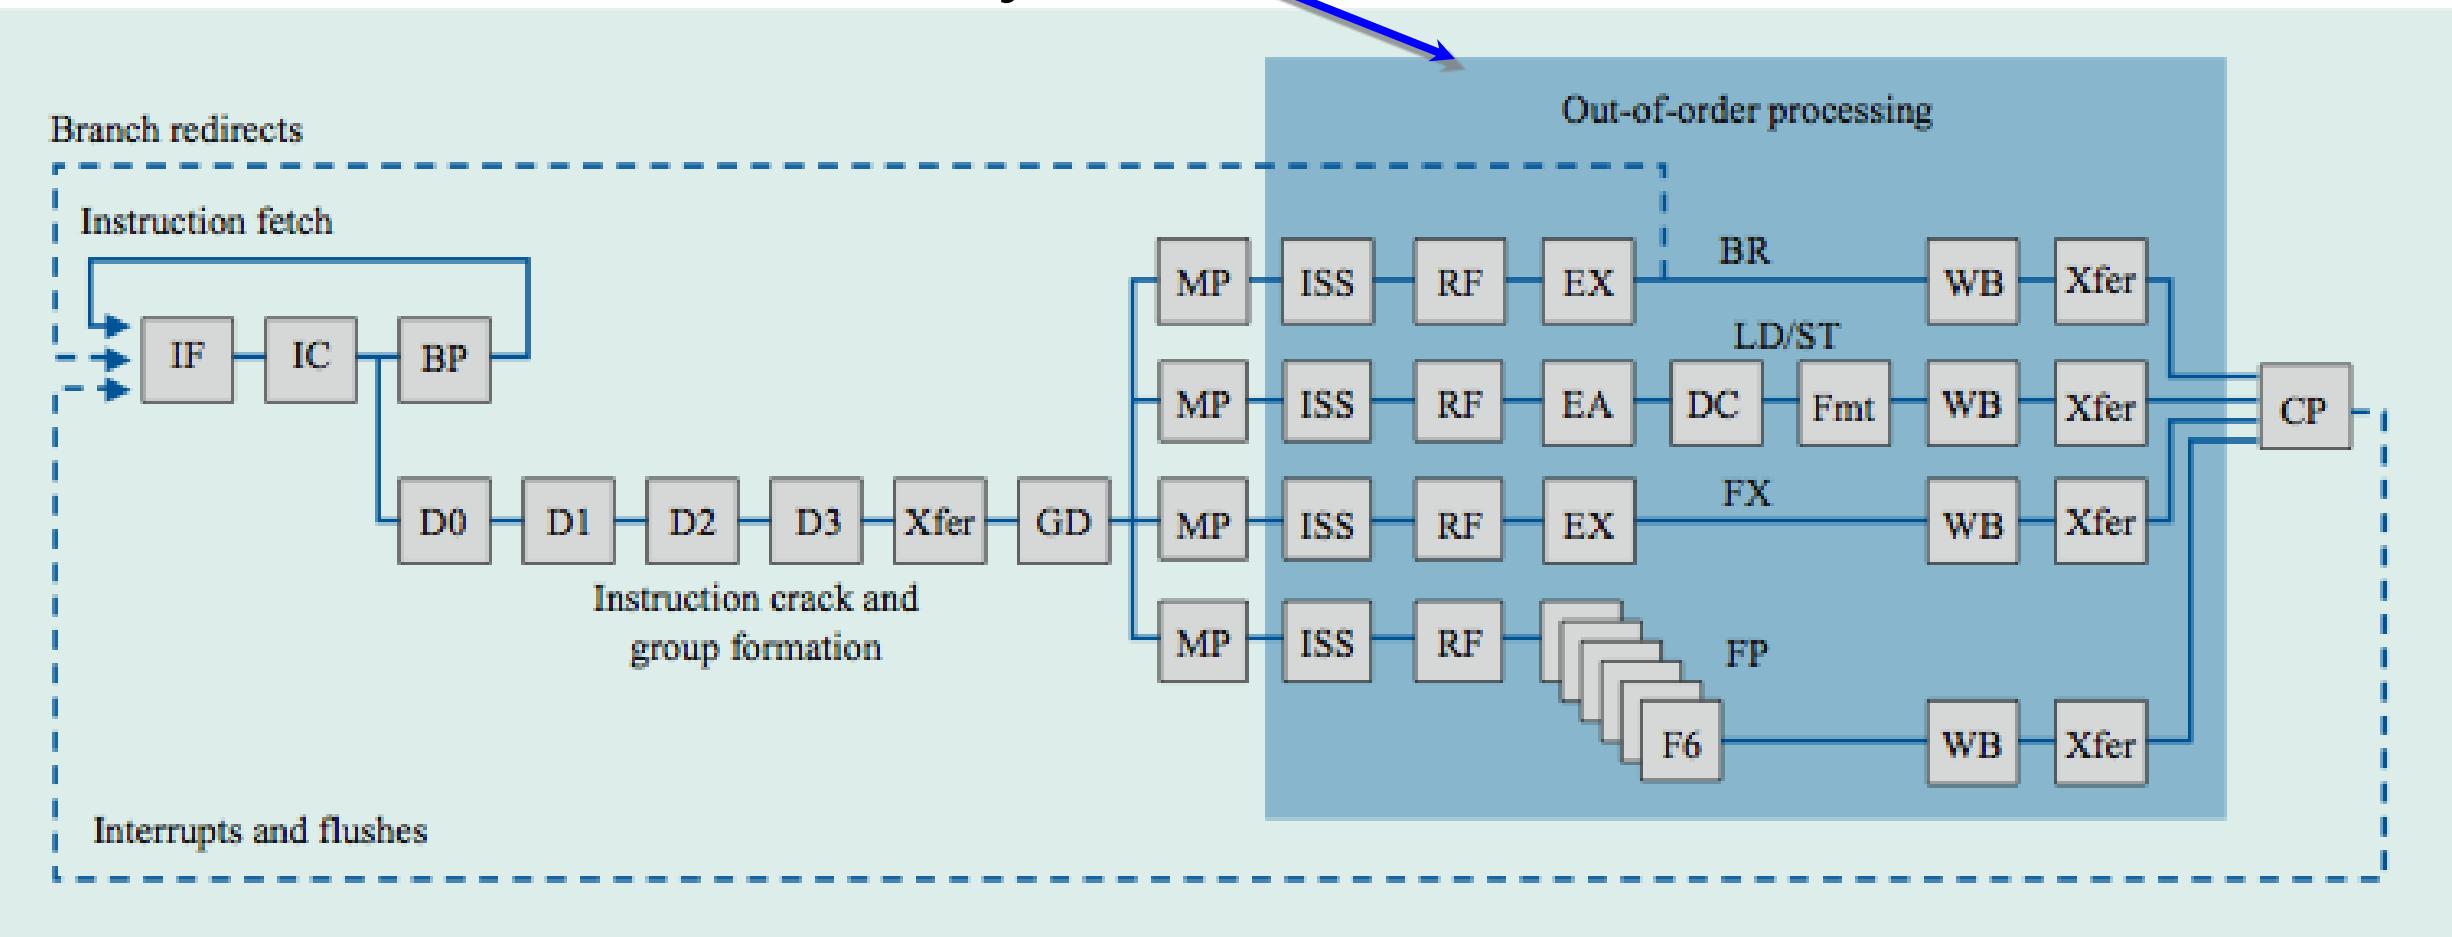
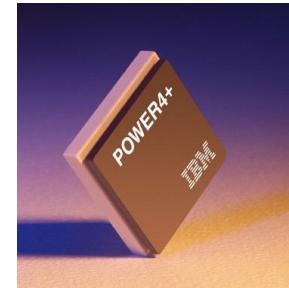
- Limitations of Instruction-Level Parallelism (ILP)
- Thread-Level Parallelism and Multithreading
- Case Studies of Multithreaded Processors
 - ◆ IBM Power 4 (one thread) versus Power 5 (two threads)
 - ◆ Fine-grained Multithreading on Sun T1 (no SMT)
 - ◆ Simultaneous Multithreading on Superscalar Processors
 - ◆ Putting It All Together: Intel Core i7 versus ARM Cortex-A53

Five Generations of IBM Processors

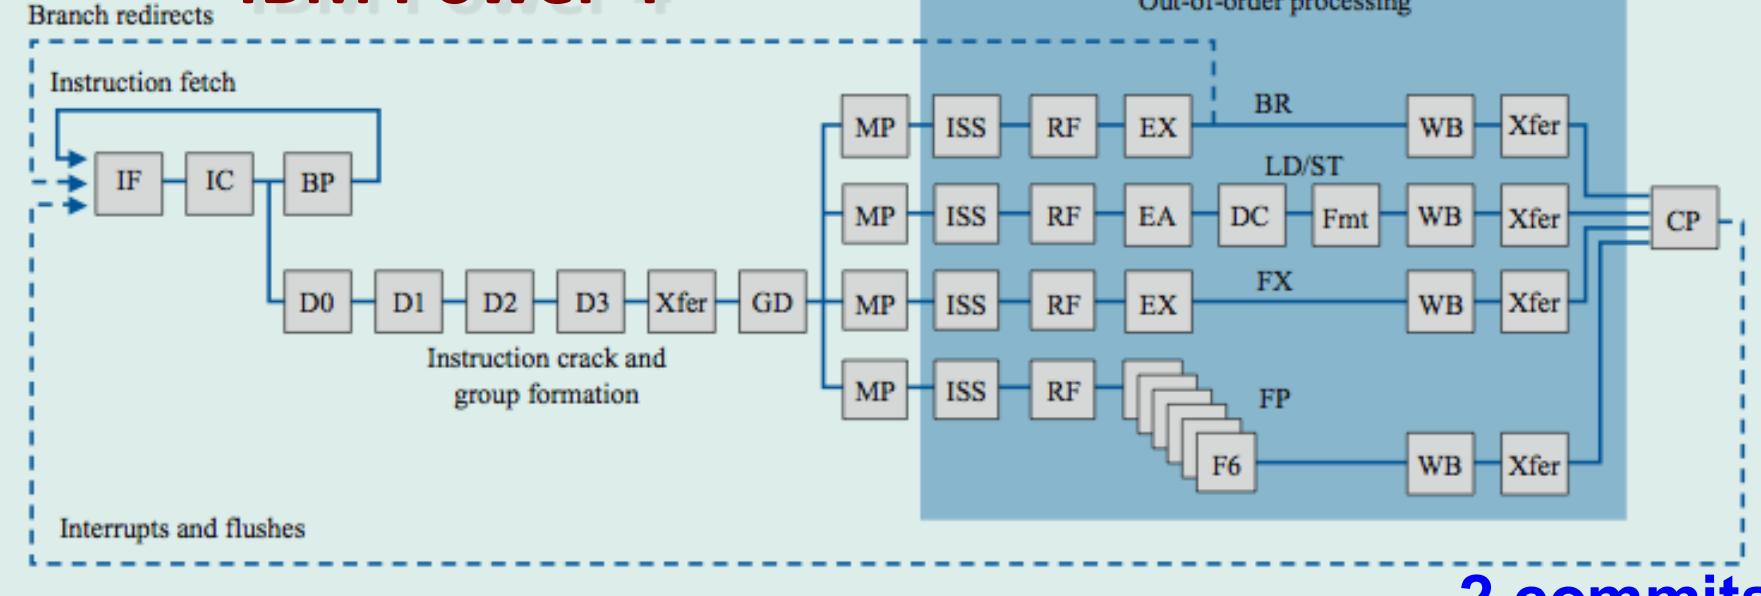
| | Power4 | Power5 | Power6 | Power7 | Power8 |
|---------------------------|--------|--------|--------|--------|---------|
| Introduced | 2001 | 2004 | 2007 | 2010 | 2014 |
| Initial clock rate (GHz) | 1.3 | 1.9 | 4.7 | 3.6 | 3.3 GHz |
| Transistor count (M) | 174 | 276 | 790 | 1200 | 4200 |
| Issues per clock | 5 | 5 | 7 | 6 | 8 |
| Functional units per core | 8 | 8 | 9 | 12 | 16 |
| SMT threads per core | 0 | 2 | 2 | 4 | 8 |
| Cores/chip | 2 | 2 | 2 | 8 | 12 |
| SMT threads per core | 0 | 2 | 2 | 4 | 8 |
| Total on-chip cache (MiB) | 1.5 | 2 | 4.1 | 32.3 | 103.0 |

IBM Power 4

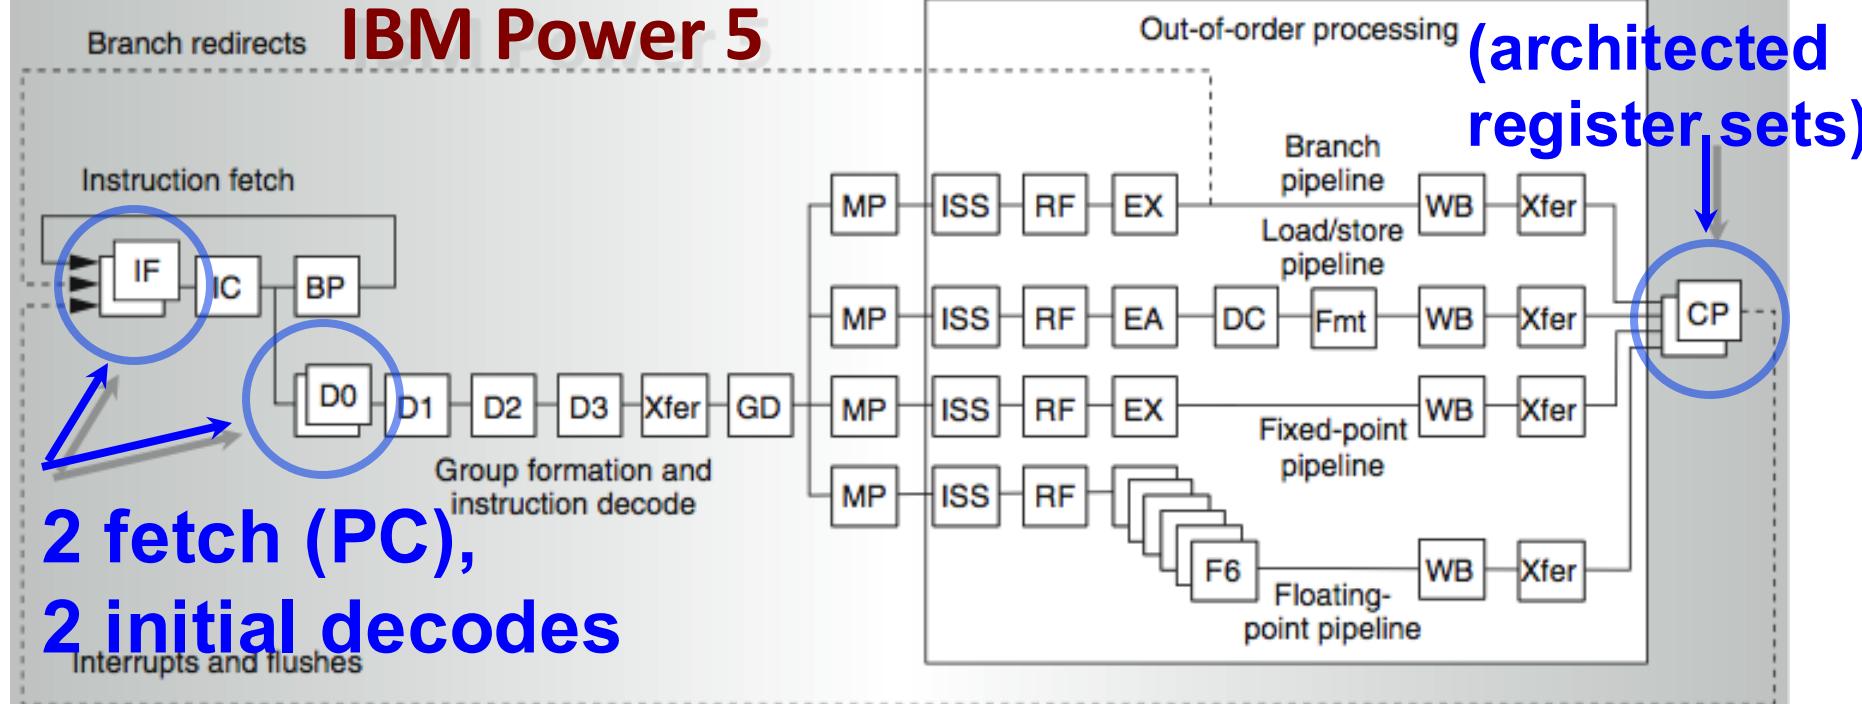
Single-threaded predecessor to Power 5. Eight execution units in out-of-order engine, each may issue an instruction each cycle.



IBM Power 4



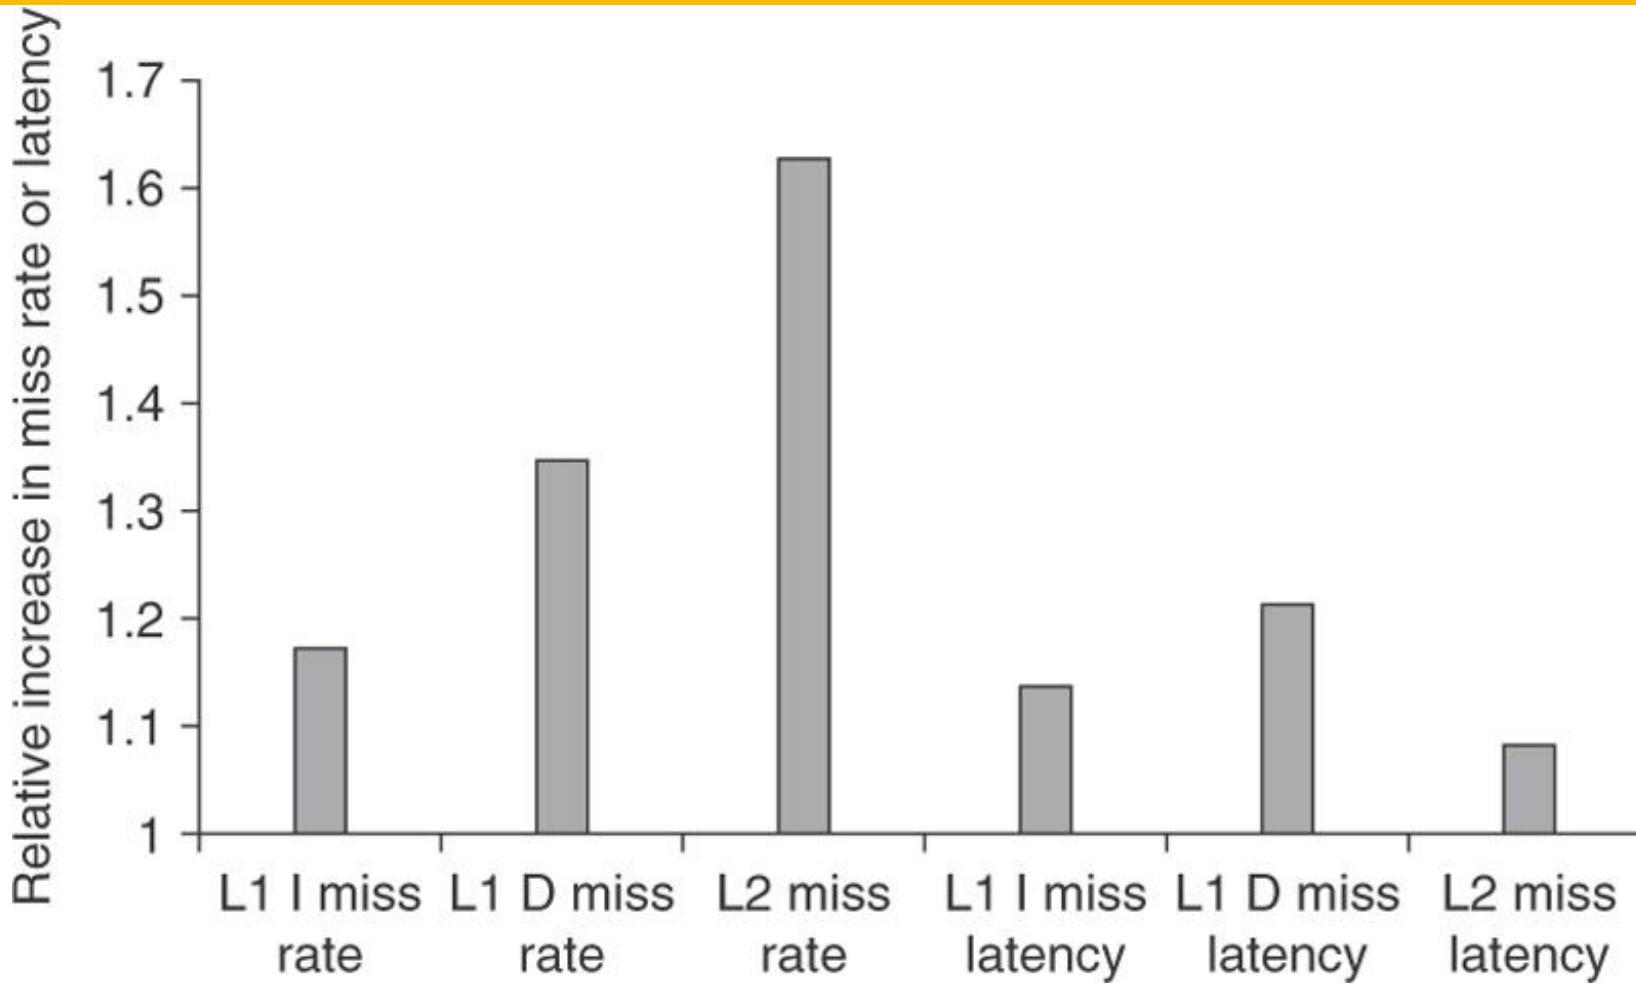
IBM Power 5



Thread-Level Parallelism and Multithreading

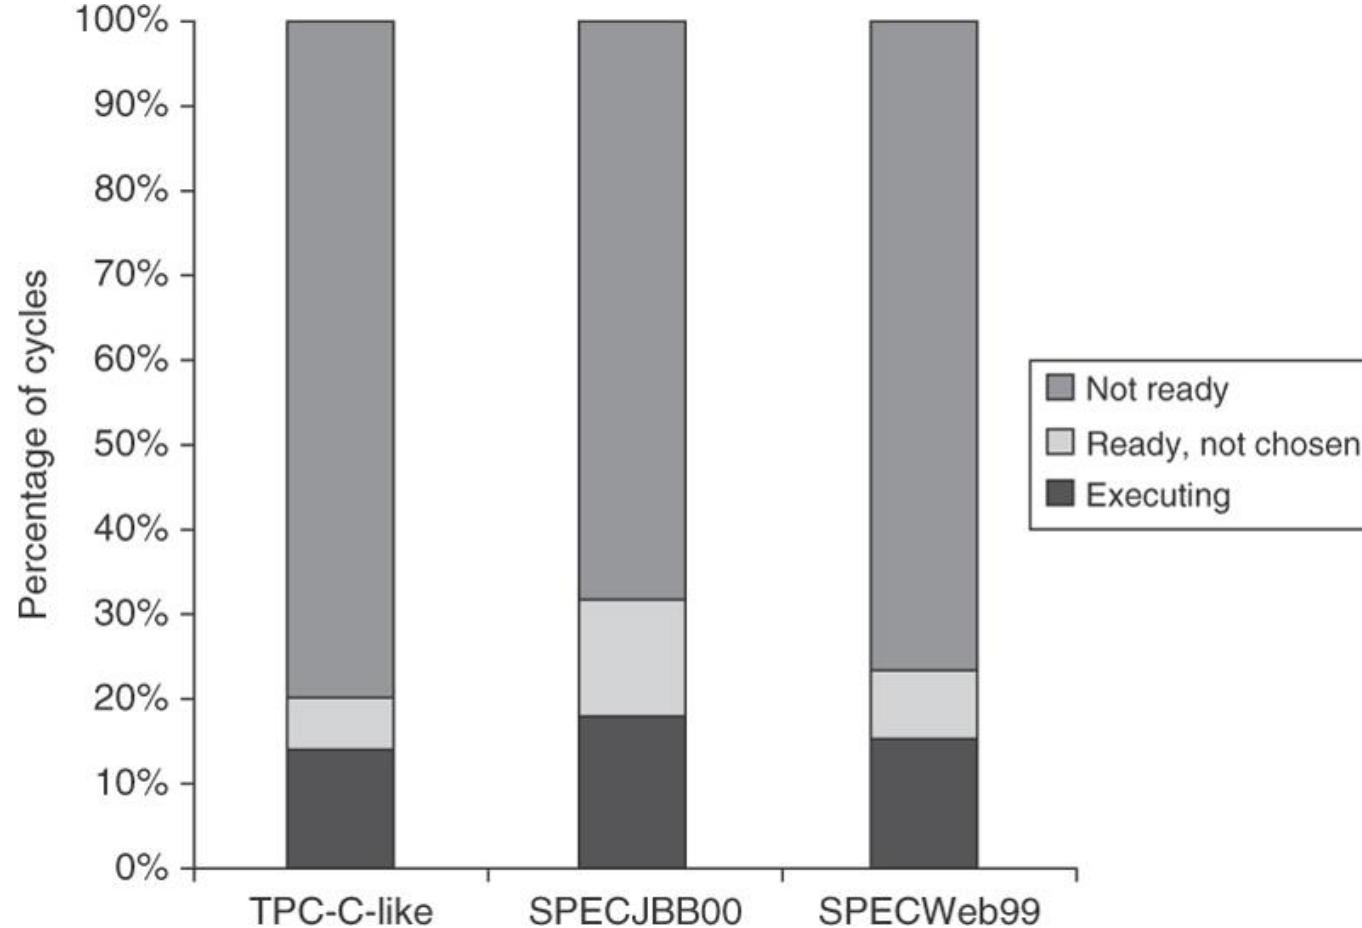
- Limitations of Instruction-Level Parallelism (ILP)
- Thread-Level Parallelism and Multithreading
- Case Studies of Multithreaded Processors
 - ◆ IBM Power 4 (one thread) versus Power 5 (two threads)
 - ◆ **Fine-grained Multithreading on Sun T1 (no SMT)**
 - ◆ Simultaneous Multithreading on Superscalar Processors
 - ◆ Putting It All Together: Intel Core i7 versus ARM Cortex-A53

Change in Miss Rates and Latencies



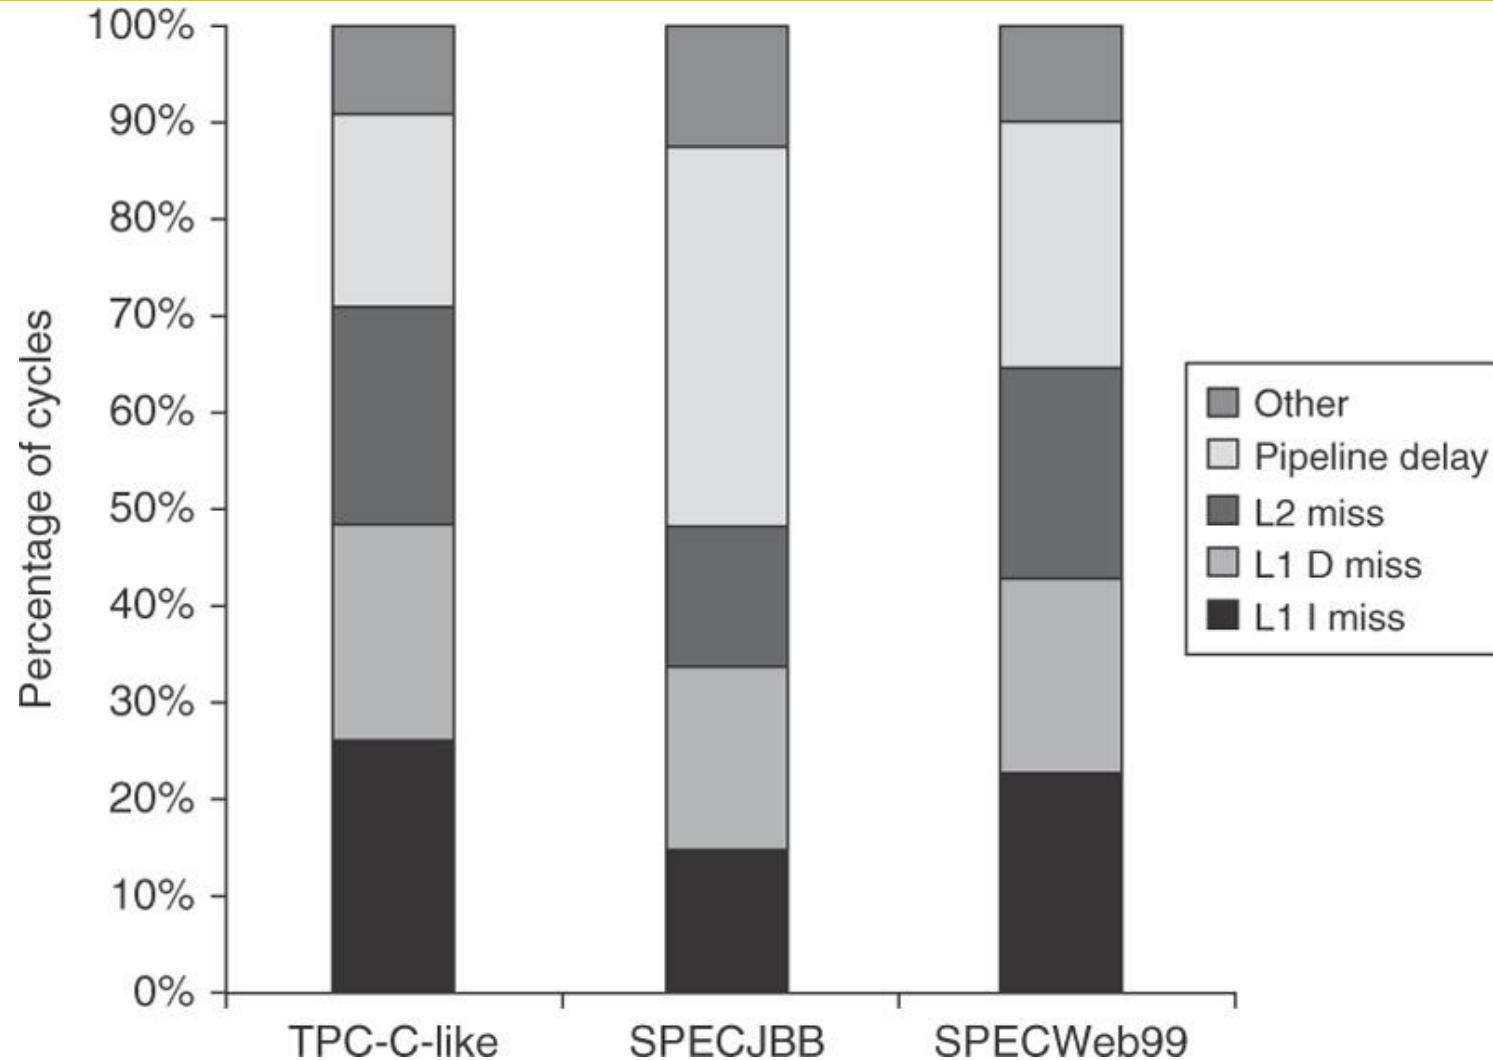
The relative change in the miss rates and miss latencies when executing with one thread per core versus four threads per core on the TPC-C benchmark on Sun UltraSPARC T1. The latencies are the actual time to return the requested data after a miss. In the four-thread case, the execution of other threads could potentially hide much of this latency.

Average Thread Statistics



“Executing” indicates the thread issues an instruction in that cycle. “Ready but not chosen” means it could issue but another thread has been chosen, and “not ready” indicates that the thread is awaiting the completion of an event. Three server benchmarks: TPC-C (online transaction processing), SPECJBB00 (SPEC Java Business Benchmark) and SPECWeb99.

Why a Thread is not Ready?

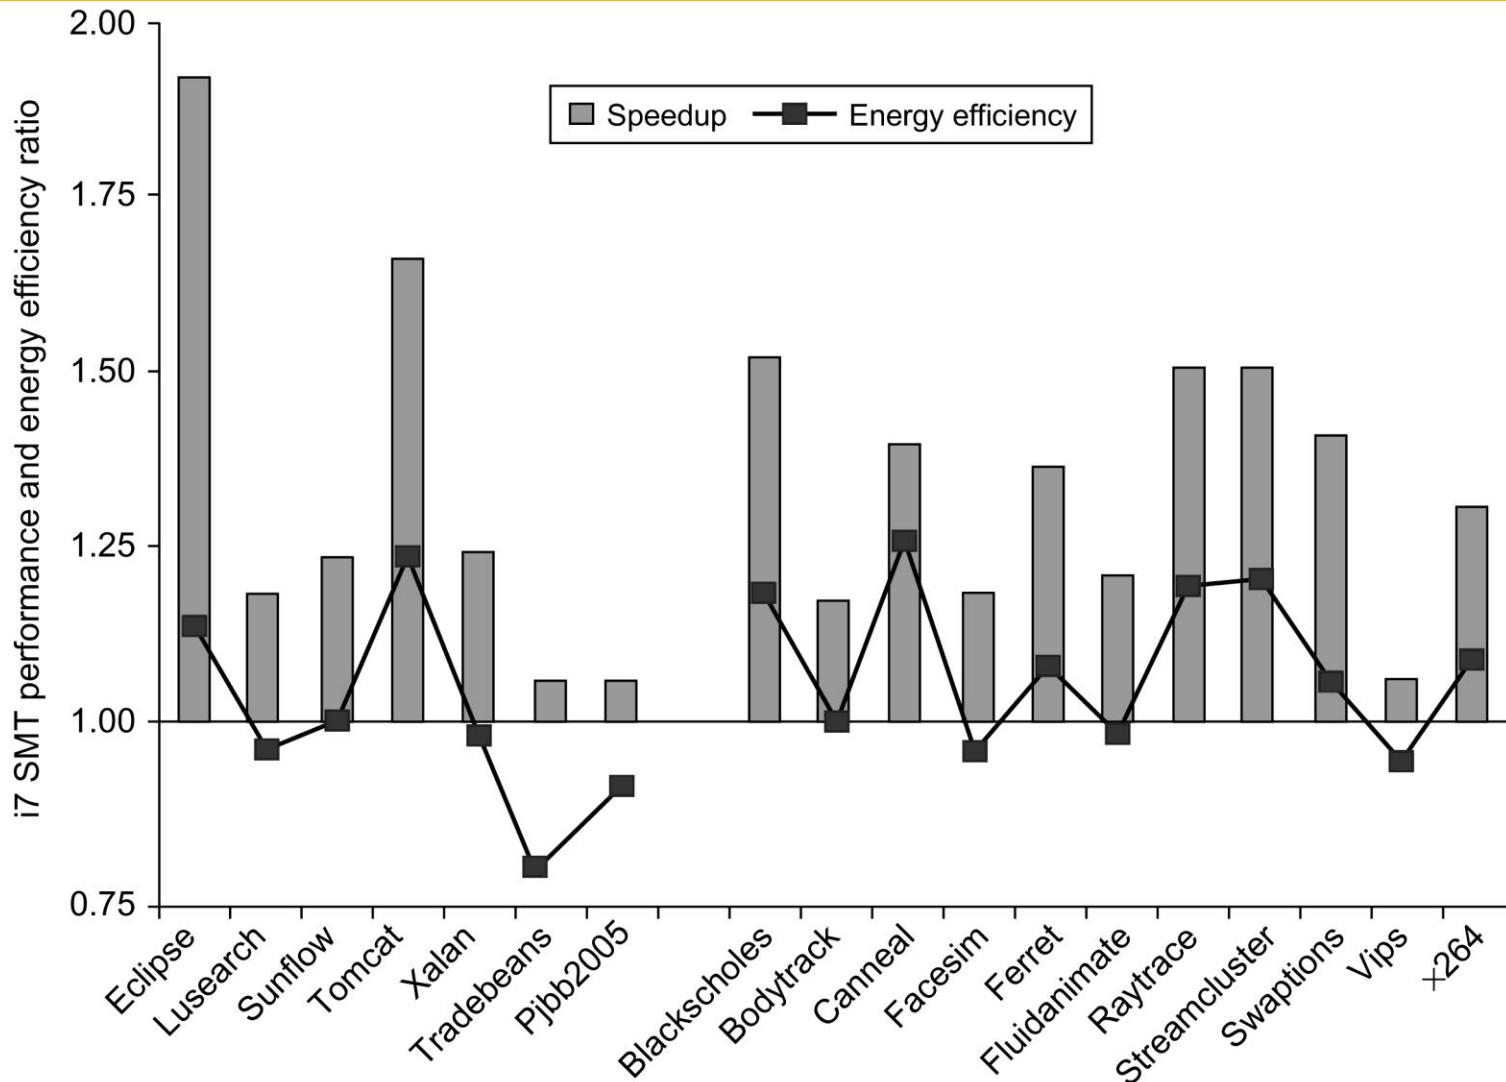


In TPC-C, store buffer full is the largest contributor (L1/L2 misses); in SPECJBB, atomic instructions are the largest contributor (pipeline delay); and in SPECWeb99, both factors contribute.

Instruction-Level Parallelism and Its Exploitation

- Limitations of Instruction-Level Parallelism (ILP)
- Thread-Level Parallelism and Multithreading
- Case Studies of Multithreaded Processors
 - ◆ IBM Power 4 (one thread) versus Power 5 (two threads)
 - ◆ Fine-grained Multithreading on Sun T1 (no SMT)
 - ◆ **Simultaneous Multithreading on Superscalar Processors**
 - ◆ Putting It All Together: Intel Core i7 versus ARM Cortex-A53

Speedup versus Energy Efficiency

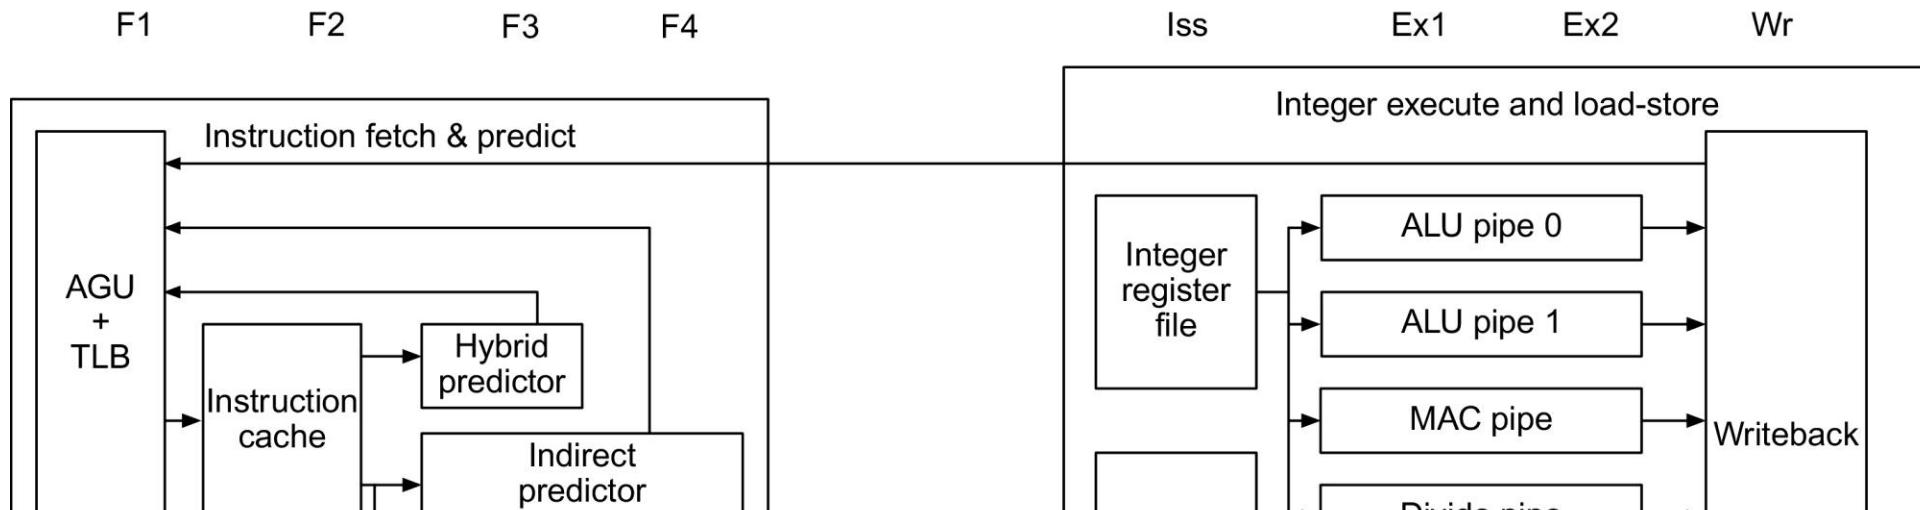


Speedup using multithreading (2 threads per core) on one core on **i7-920** averages 1.28 for the Java benchmarks and 1.31 for the PARSEC benchmarks. The energy efficiency averages 0.99 & 1.07, respectively. Anything above 1.0 for energy efficiency indicates that the feature reduces execution time by more than it increases average power.

Thread-Level Parallelism and Multithreading

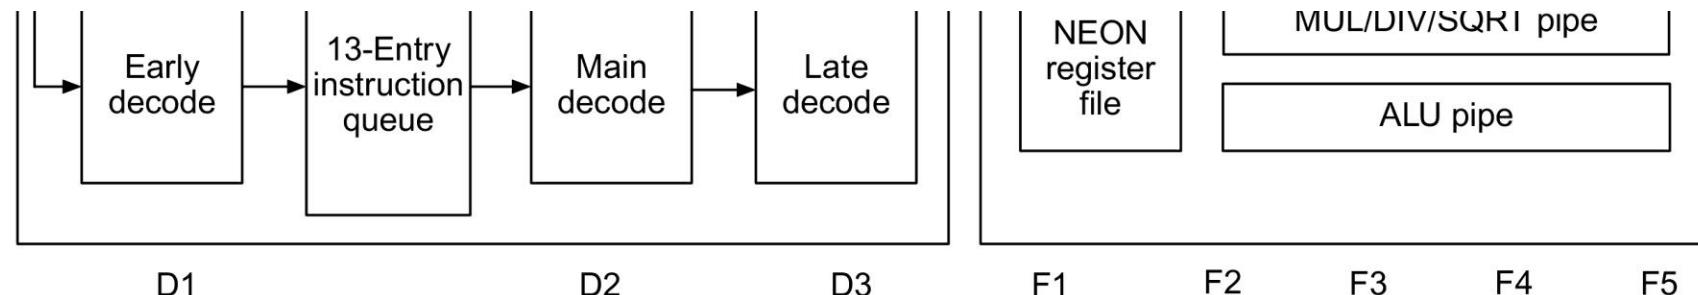
- Limitations of Instruction-Level Parallelism (ILP)
- Thread-Level Parallelism and Multithreading
- Case Studies of Multithreaded Processors
 - ◆ IBM Power 4 (one thread) versus Power 5 (two threads)
 - ◆ Fine-grained Multithreading on Sun T1 (no SMT)
 - ◆ Simultaneous Multithreading on Superscalar Processors
- Putting It All Together: Intel i7 versus ARM Cortex-A53

ARM Cortex-A53

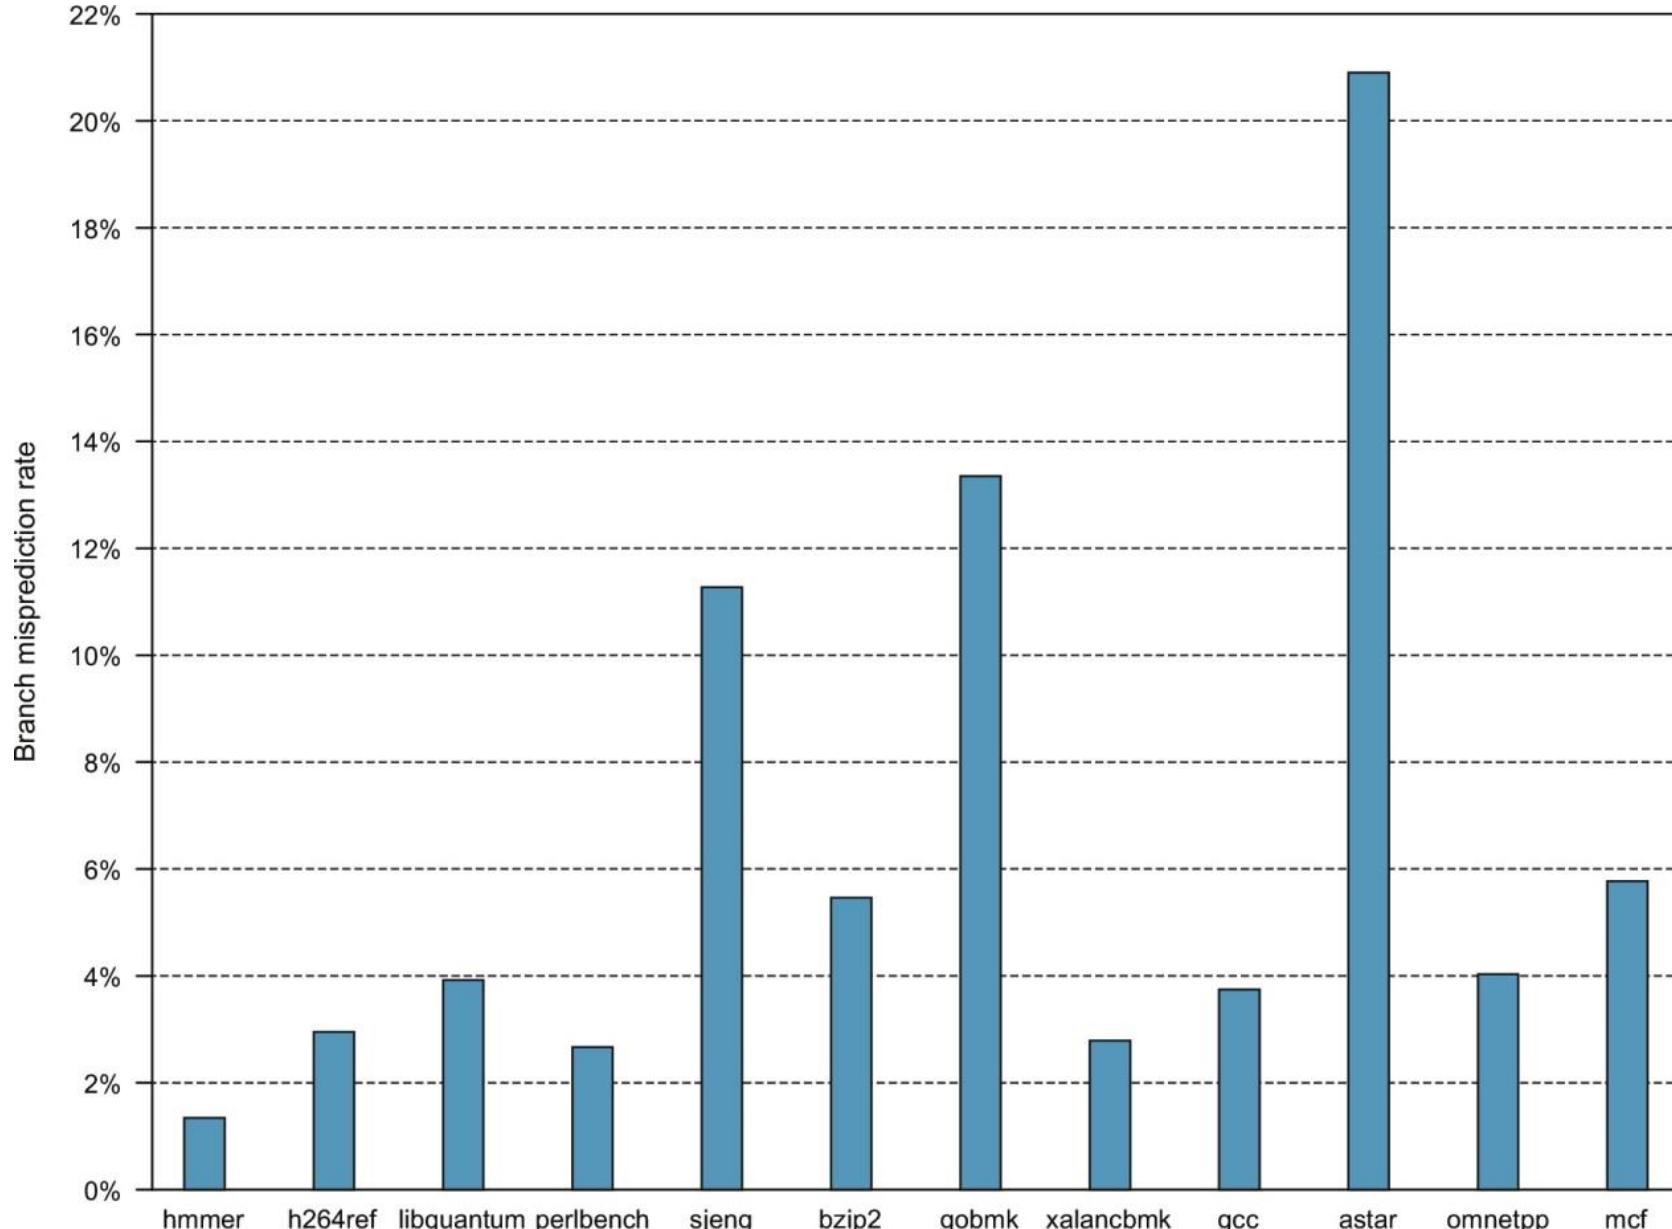


A53 integer pipeline has 8 stages and floating-point pipeline is 10 stages

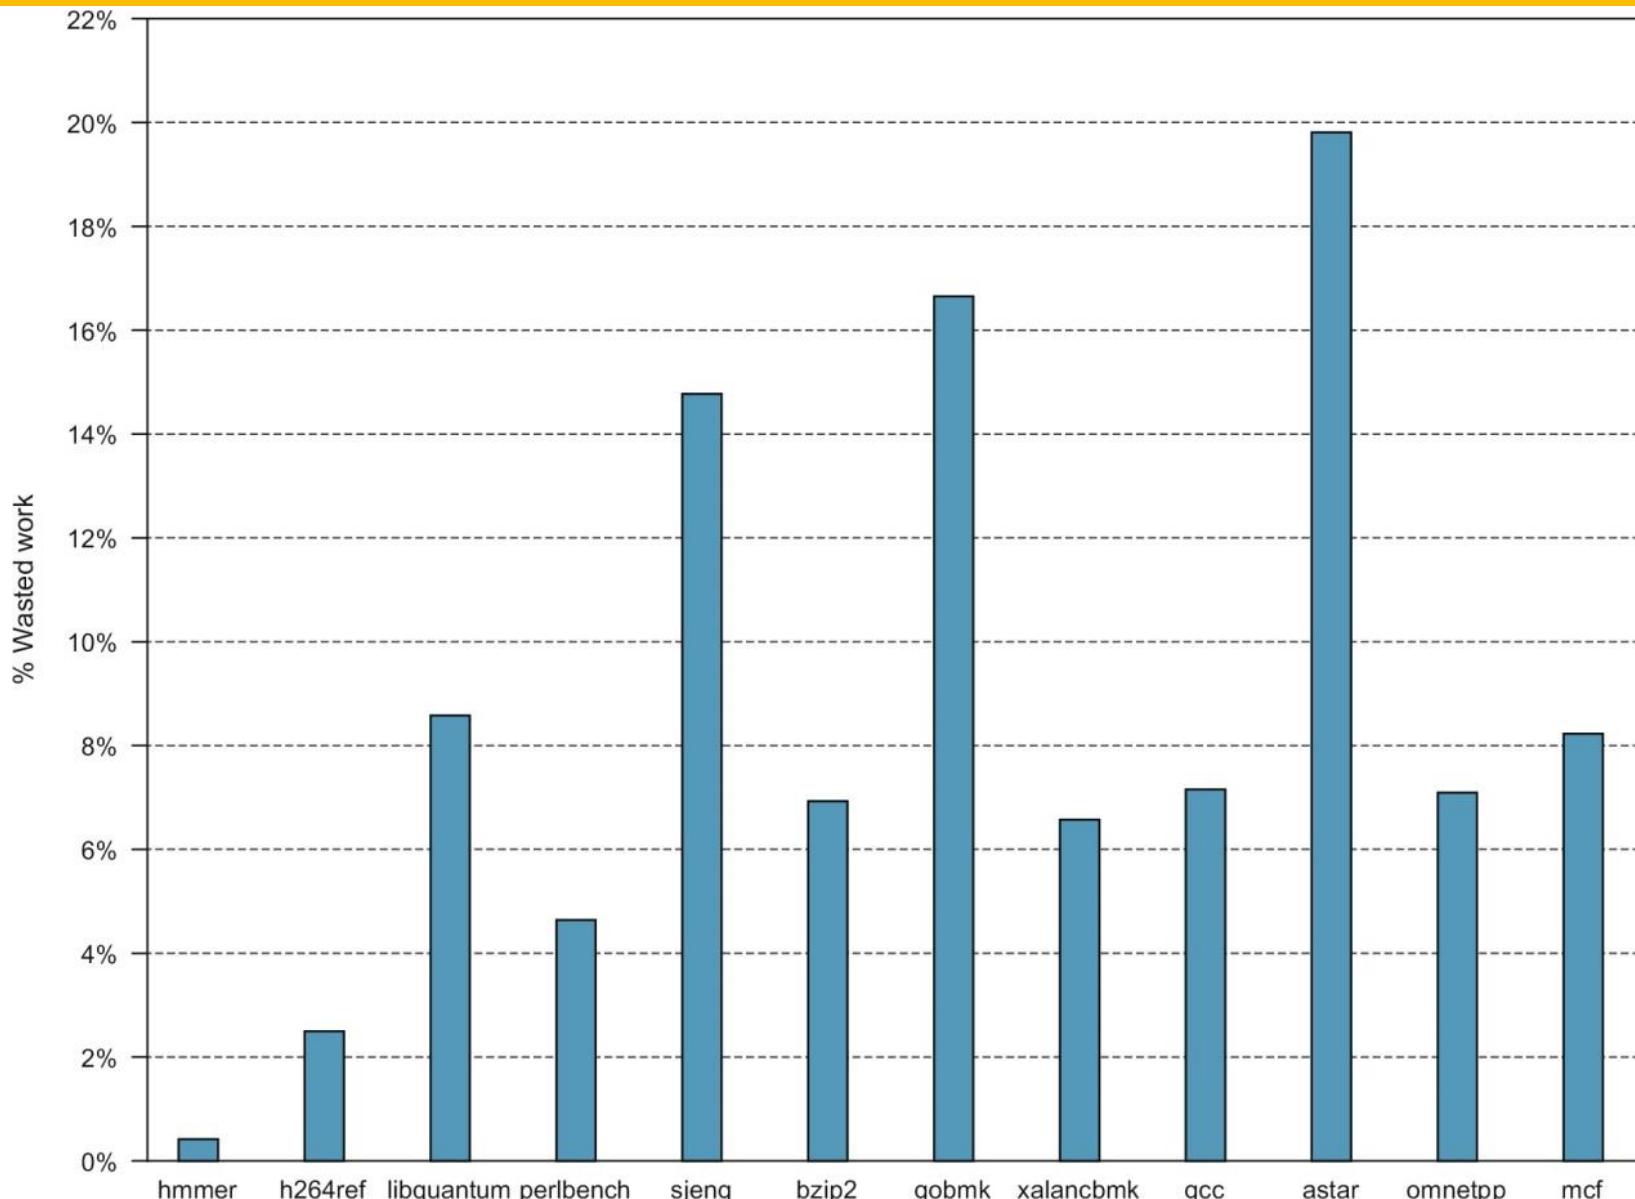
- F1 and F2 fetch the instruction
- D1 and D2 perform basic decoding, D3 for decoding complex instructions
- Issue (Iss), Ex1, Ex2, and Wr complete the integer pipeline
- Floating-point execution pipeline is 5 stages + 5 cycles for fetch and decode (10 stages in total)



Misprediction Rate in A53 for SPECint2006

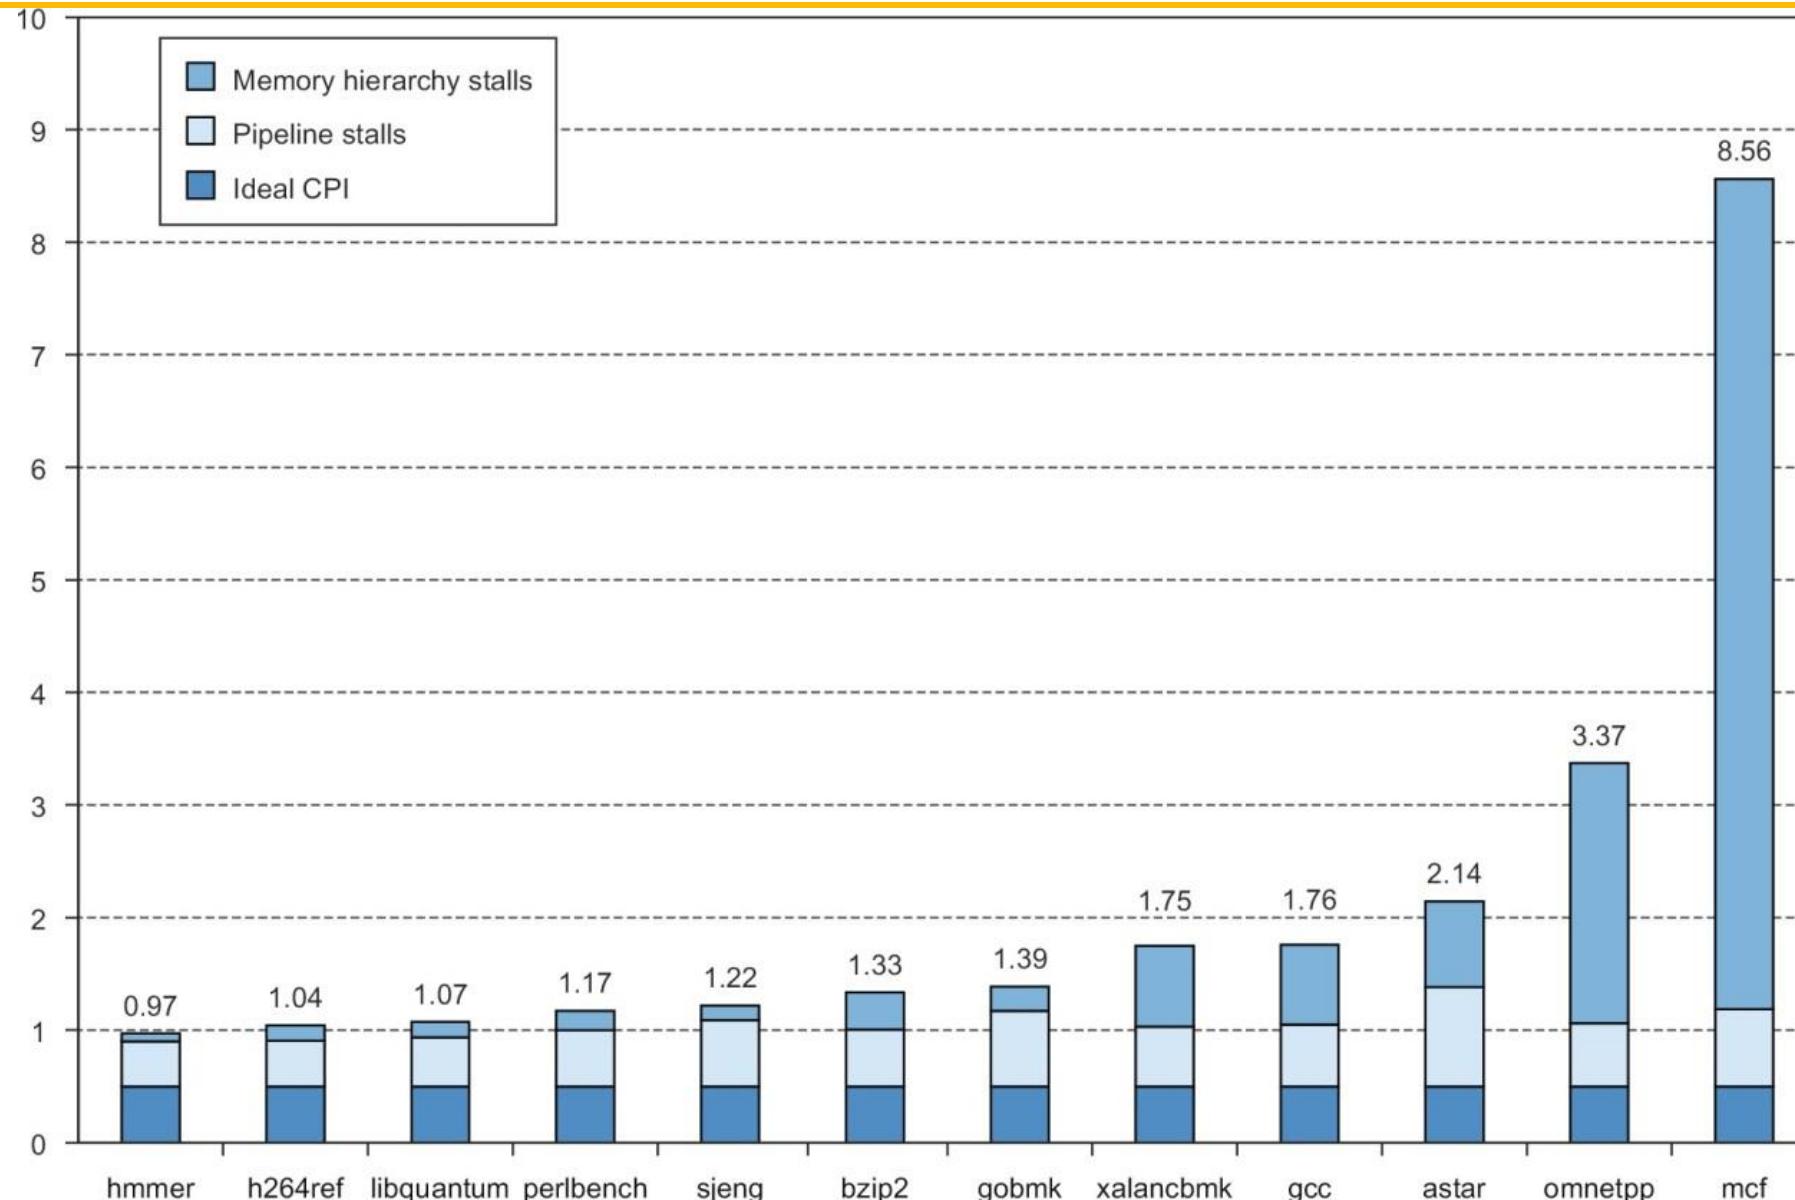


Wasted Work in A53



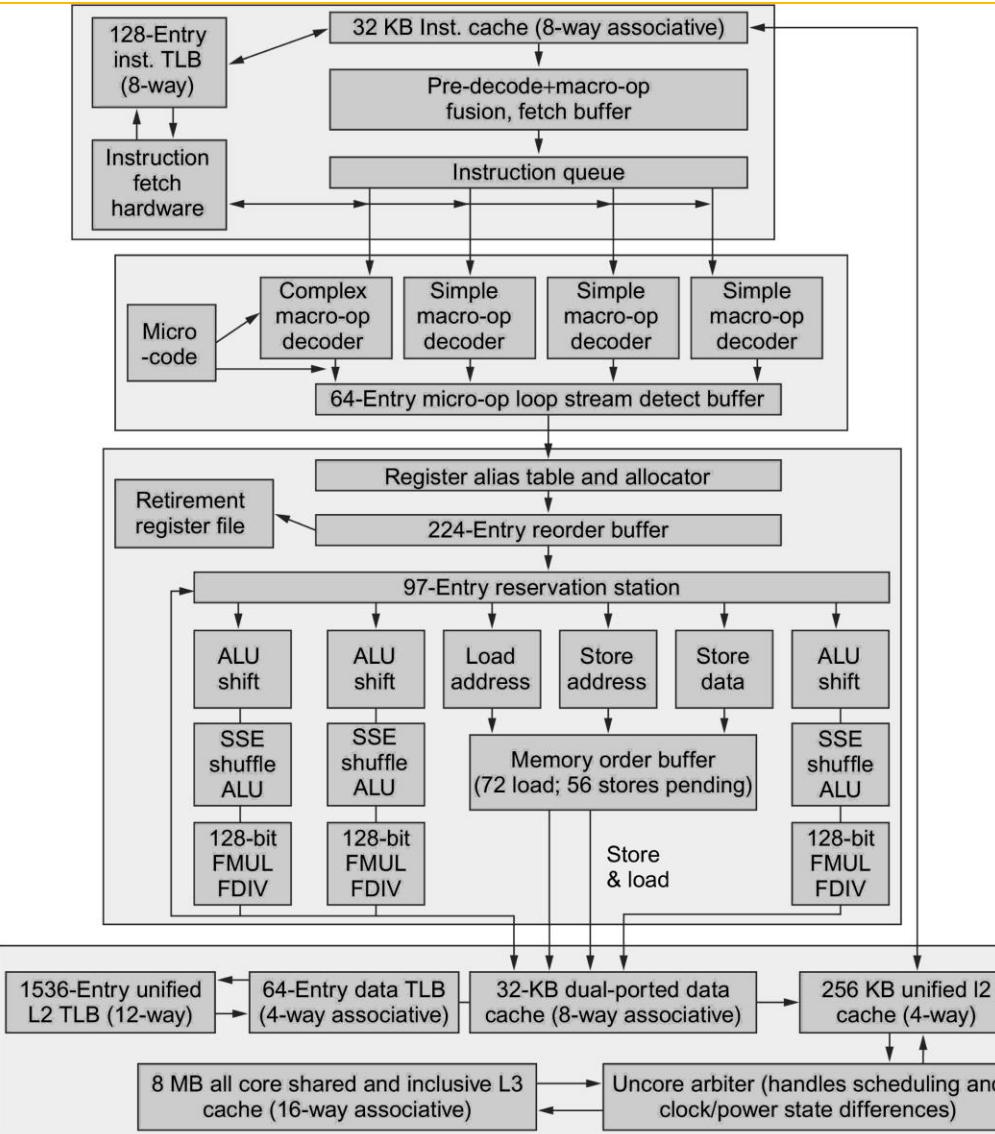
The wasted effort depends on various factors, including branch misprediction, data dependence, and cache misses.

CPI in A53



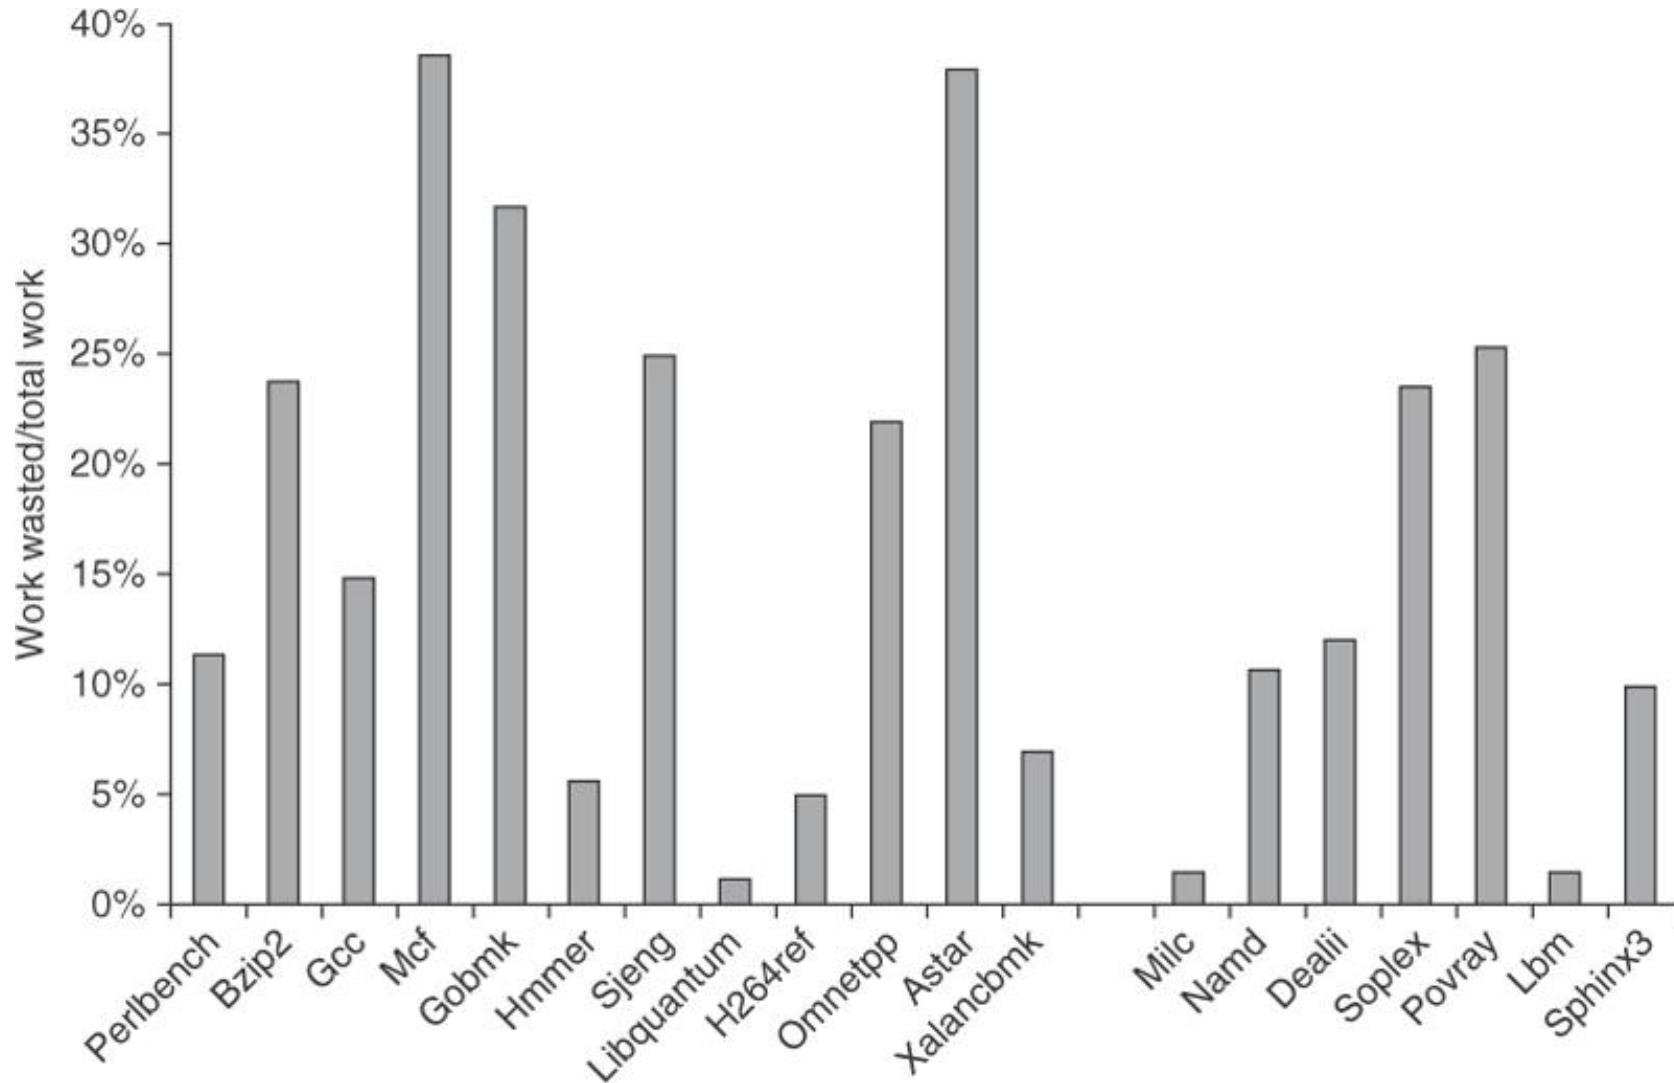
Ideal CPI is 0.5. Pipeline stalls are significant, but outweighed by cache misses in few cases.

Intel Core i7



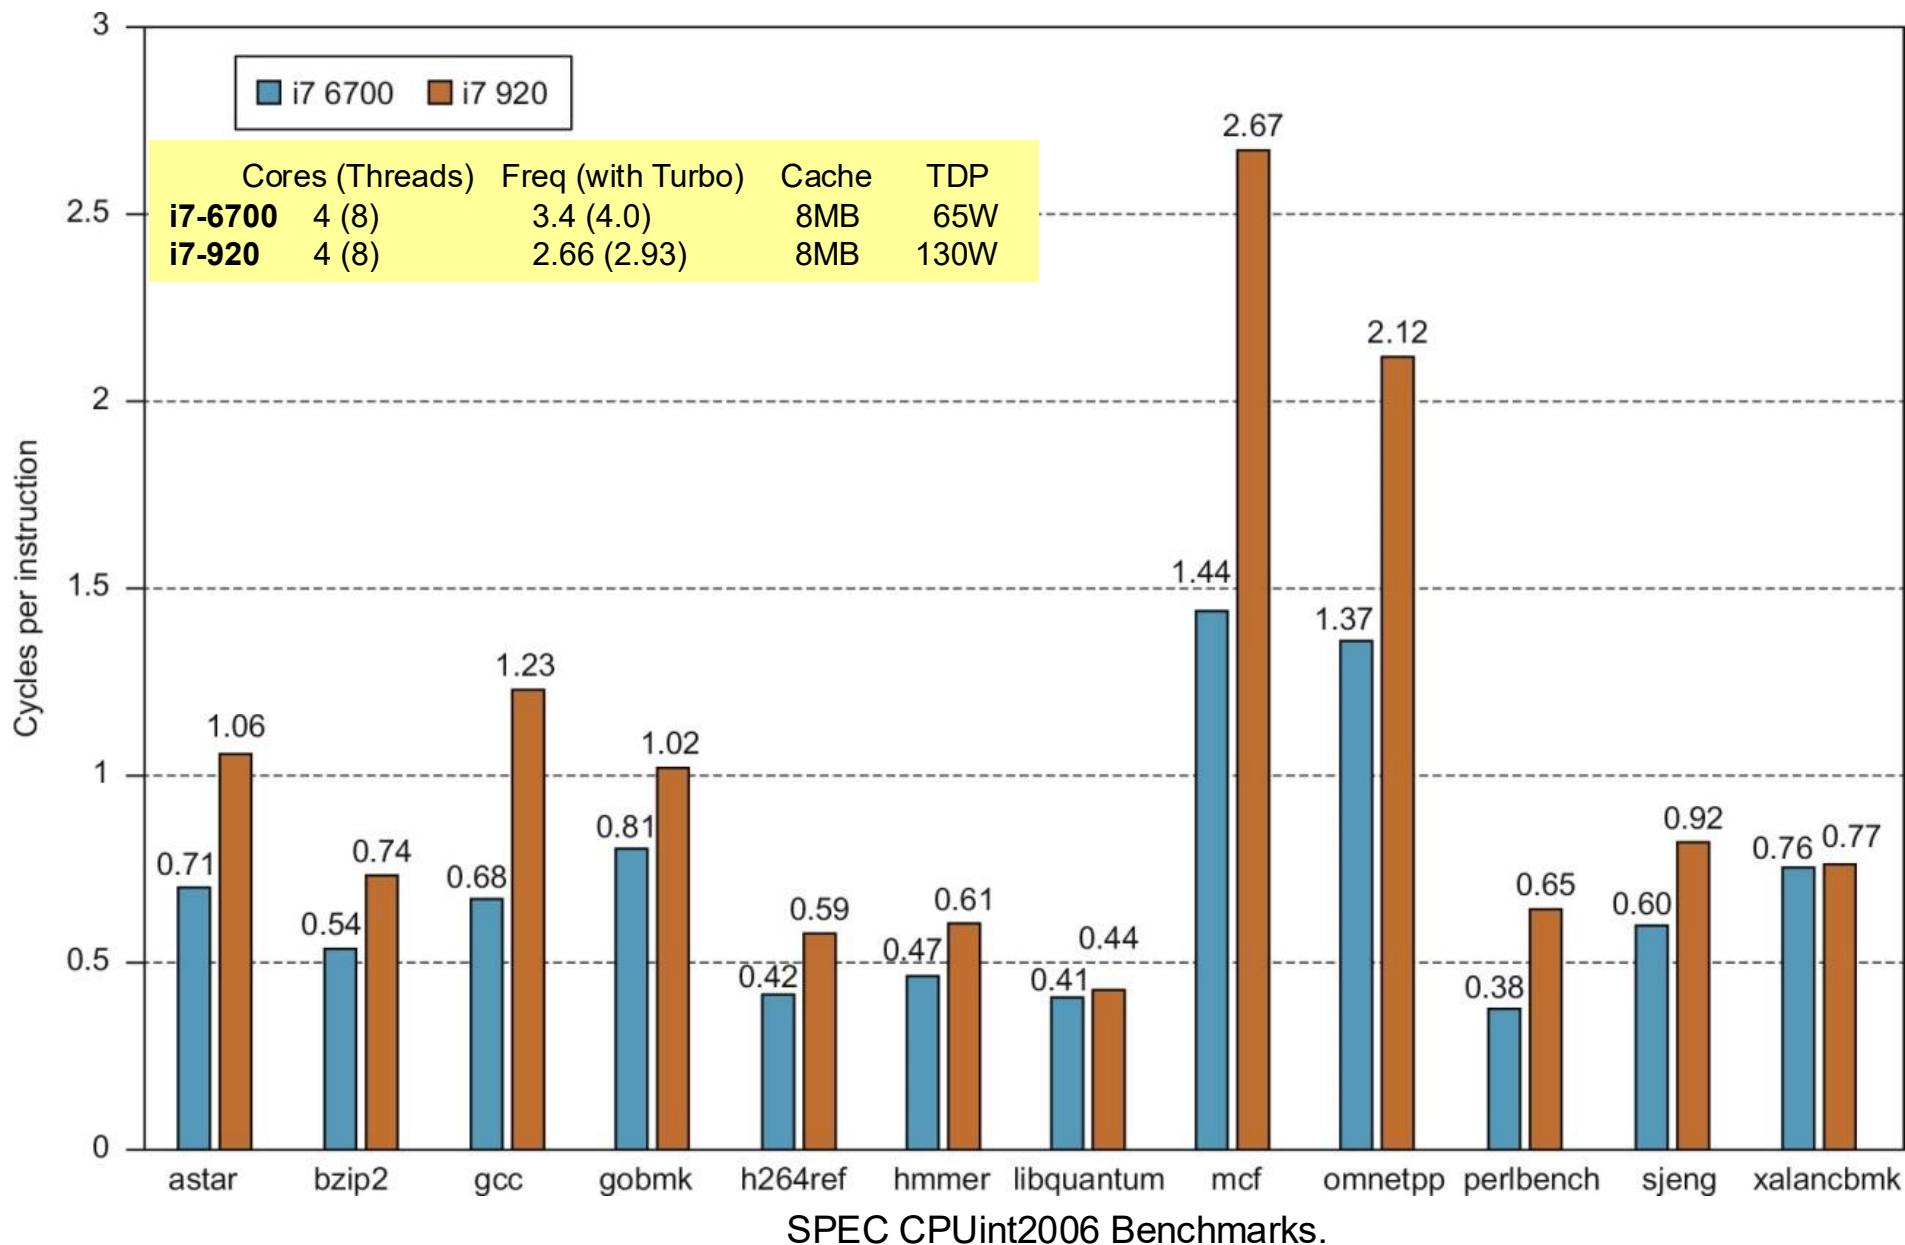
Pipeline depth is 14 stages, with branch mispredictions costing 17 cycles. There are 48 load and 32 store buffers. Six independent functional units can each begin execution of a ready micro-op in the same cycle.

Wasted Micro-Ops in i7



The amount of "**wasted work**" is plotted by taking the ratio of dispatched micro-ops that do not graduate to all dispatched micro-ops. For example, the ratio is 25% for *sjeng*, meaning that 25% of the dispatched and executed micro-ops are thrown away.

CPI of i7 6700 and i7 920



Intel i7 920 versus i7 6700

| Resource | i7 920 (Nehalem) | i7 6700 (Skylake) |
|-----------------------------|------------------|-------------------|
| Micro-op queue (per thread) | 28 | 64 |
| Reservation stations | 36 | 97 |
| Integer registers | NA | 180 |
| FP registers | NA | 168 |
| Outstanding load buffer | 48 | 72 |
| Outstanding store buffer | 32 | 56 |
| Reorder buffer | 128 | 256 |

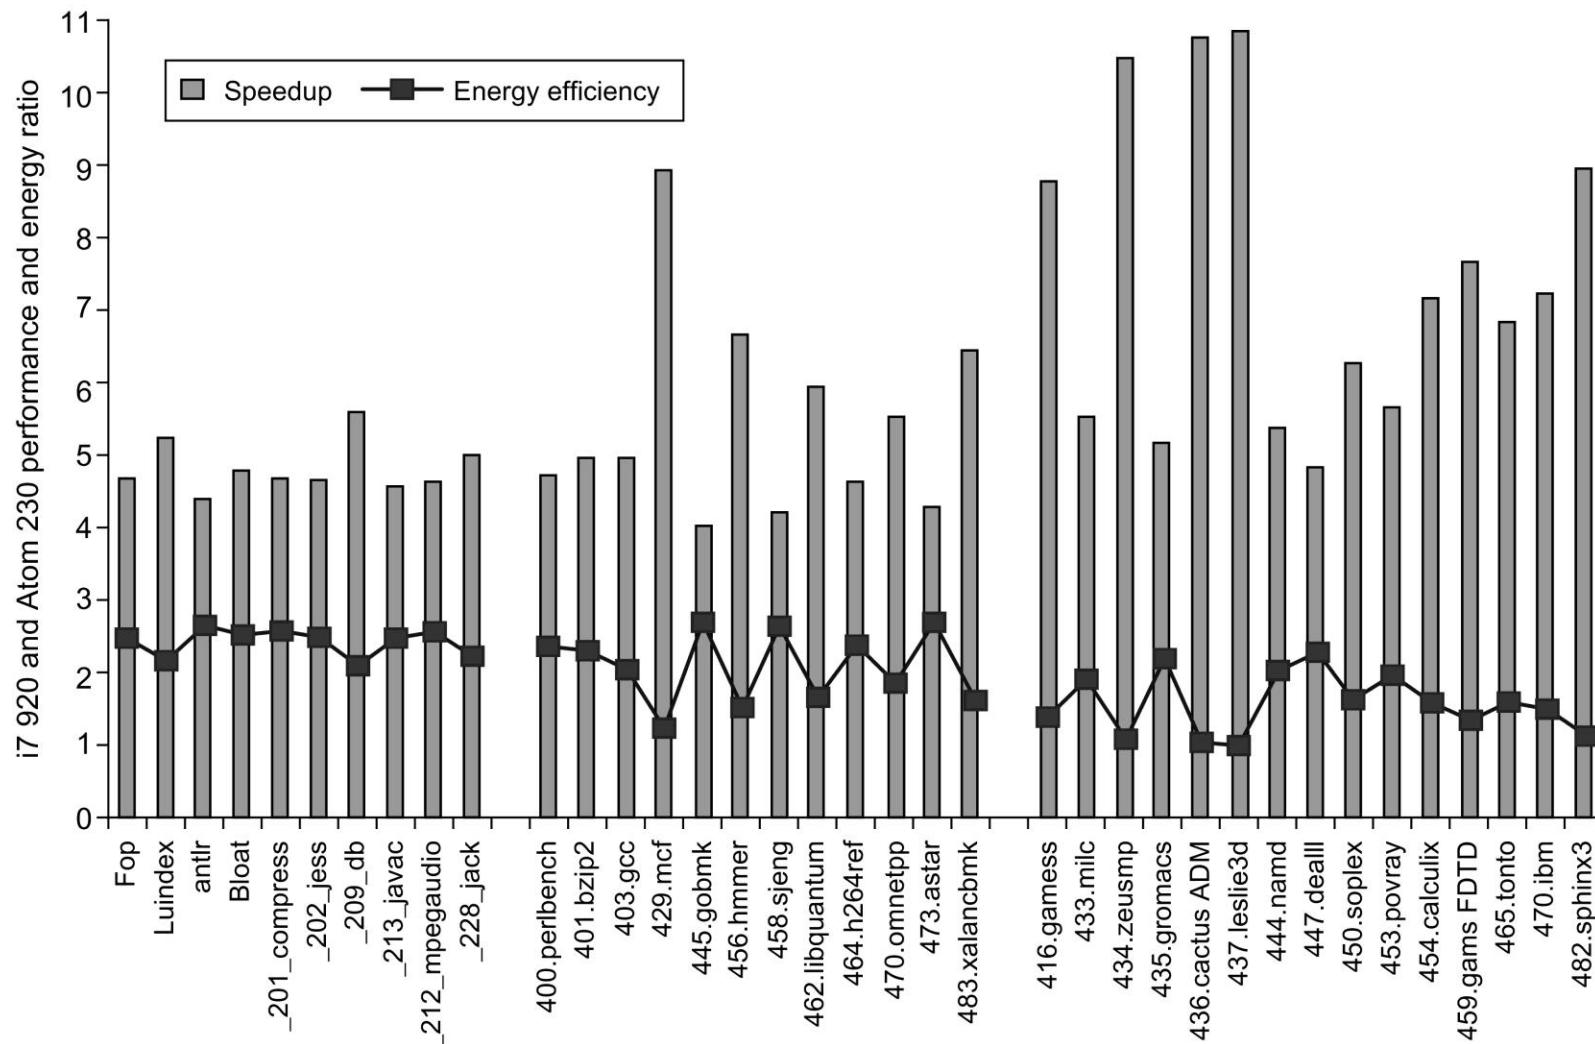
Nehalem (i7 920) used a reservation station plus reorder buffer organization. In later microarchitectures, the reservation stations serve as scheduling resources, and register renaming is used rather than the reorder buffer. The reorder buffer in the **Skylake** (i7 6700) microarchitecture serves only to buffer control information. The choices of the size of various buffers and renaming registers, while appearing sometimes arbitrary, are likely based on extensive simulation.

Intel i7 vs ARM A8 vs Atom 230

| | | Intel i7 920 | ARM A8 | Intel Atom 230 |
|--------------------------|-------------------------|---|--|---|
| Area | Specific characteristic | Four cores, each with FP | One core, no FP | One core, with FP |
| Physical chip properties | Clock rate | 2.66 GHz | 1 GHz | 1.66 GHz |
| | Thermal design power | 130 W | 2 W | 4 W |
| | Package | 1366-pin BGA | 522-pin BGA | 437-pin BGA |
| Memory system | TLB | Two-level All four-way set associative 128 I/64 D 512 L2 | One-level fully associative 32 I/32 D | Two-level All four-way set associative 16 I/16 D 64 L2 |
| | Caches | Three-level 32 KiB/32 KiB 256 KiB 2–8 MiB | Two-level 16/16 or 32/32 KiB 128 KiB–1 MiB | Two-level 32/24 KiB 512 KiB |
| | Peak memory BW | 17 GB/s | 12 GB/sec | 8 GB/s |
| Pipeline structure | Peak issue rate | 4 ops/clock with fusion | 2 ops/clock | 2 ops/clock |
| | Pipe line scheduling | Speculating out of order | In-order dynamic issue | In-order dynamic issue |
| | Branch prediction | Two-level | Two-level 512-entry BTB 4 K global history 8-entry return stack | Two-level |

Design philosophy: i7 (desktop or servers), low-power mobile devices (A8) and notebook (Atom)

Performance vs. Energy Efficiency



i7 920 is 4 to over 10 times faster than the Atom 230 but that it is about 2 times less power efficient on average! One core is active on the i7 (single threaded), and the rest are in deep power saving mode. Turbo Boost is used on the i7, which increases its performance advantage but slightly decreases its relative energy efficiency.

Conclusion

- Instruction-level parallelism (ILP) has its limits
- Exploiting thread-level parallelism to go beyond ILP
 - ◆ Coarse-grained Multithreading
 - ◆ Fine-grained Multithreading
 - ◆ Simultaneous Multithreading
- Case studies of multithreaded processors
- Please read Section 3.10 – 3.13
- We review memory hierarchy next.